

# **CS310 Operating Systems**

## **Lecture 17: Inter Process Communication - Message Queues, Pipe**

Ravi Mittal  
IIT Goa

# Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
- <https://www.softprayog.in/>

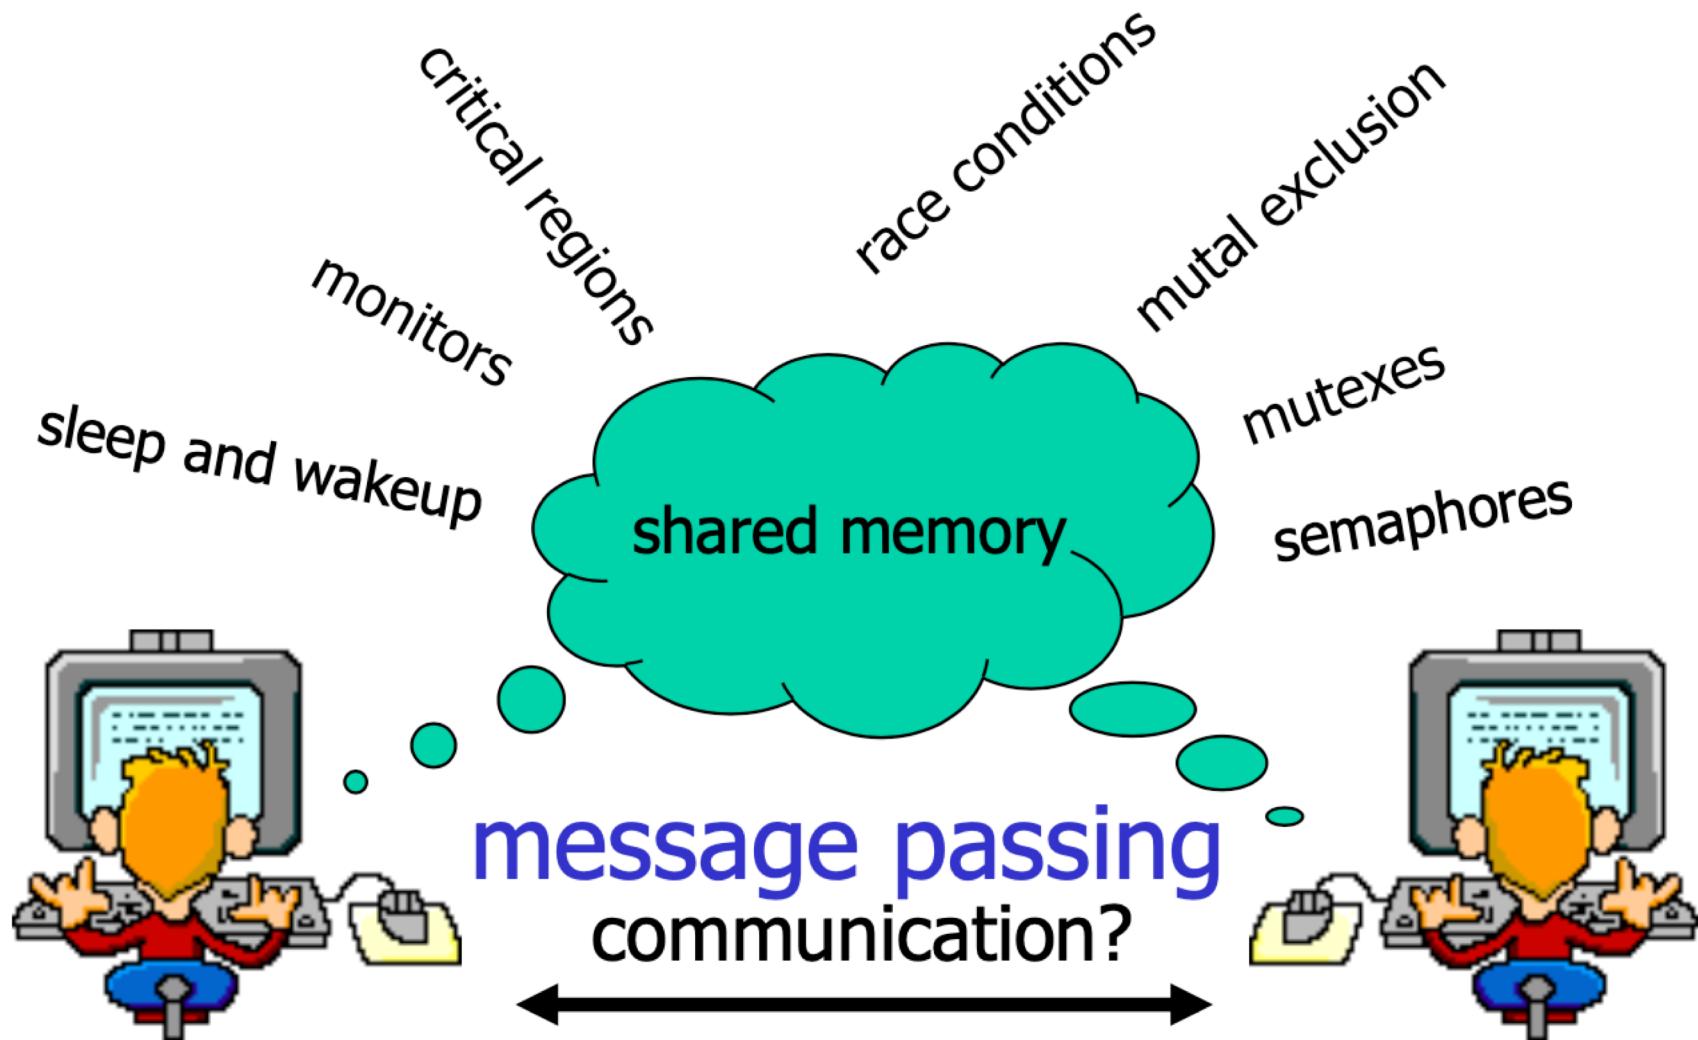
# Reading

- Book: Operating System Concepts, 10<sup>th</sup> Edition, by Silberschatz, Galvin, and Gagne
- Book: Advanced Unix Programming, by Mark J Rochkind
- Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Part 1 and Part 2

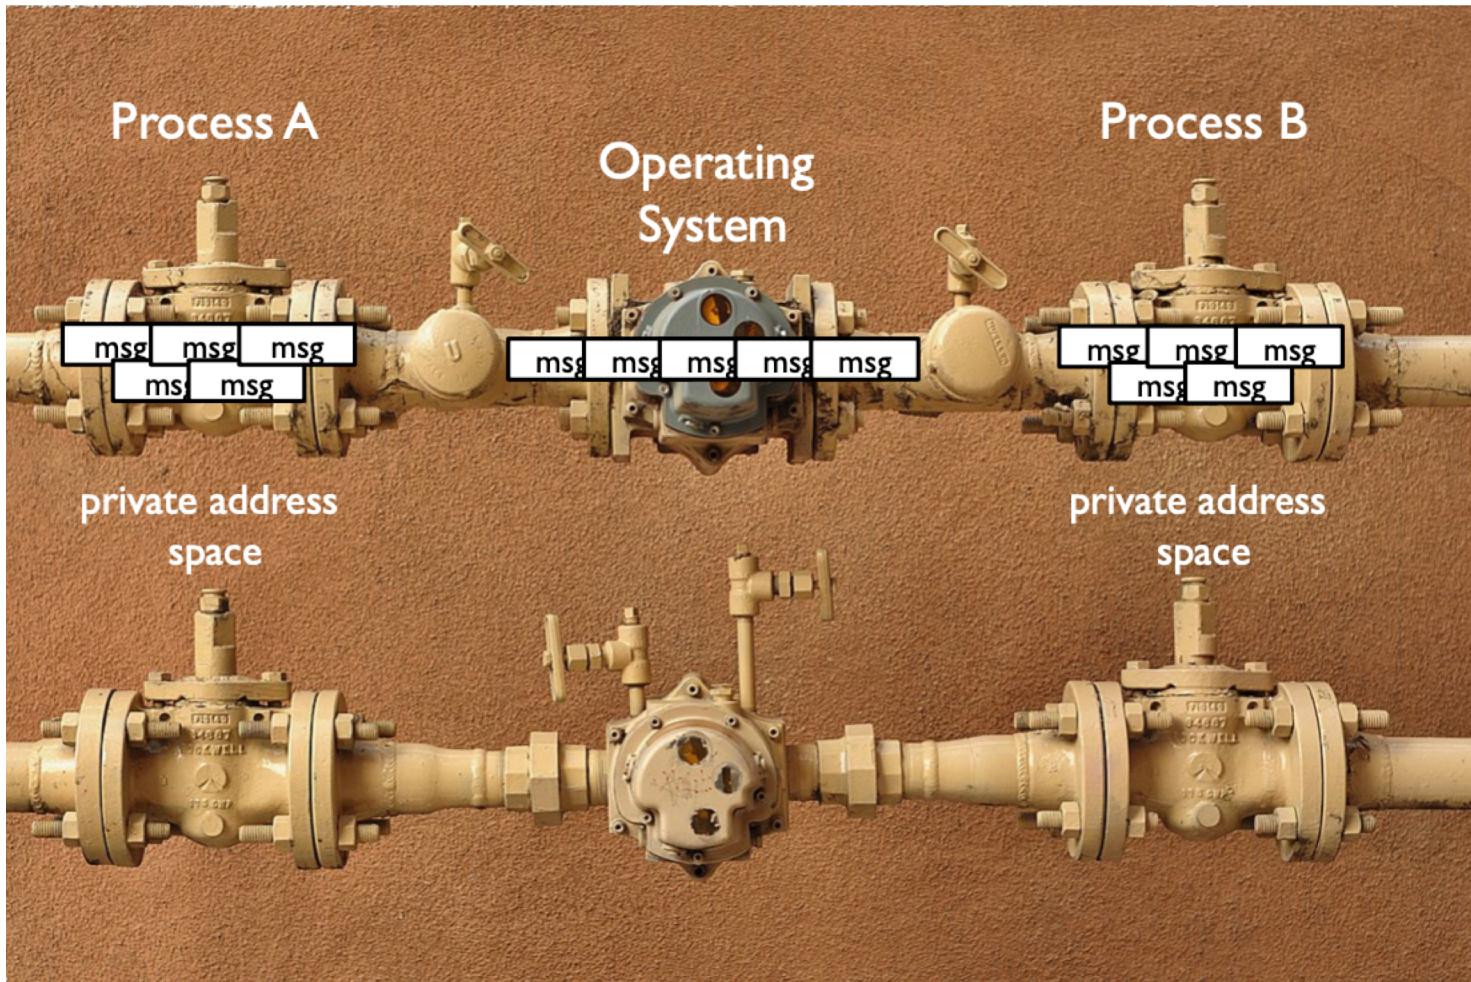
# We will study today

- Message Queues
- Pipes
- FIFO (self reading)

# IPC – Big Picture



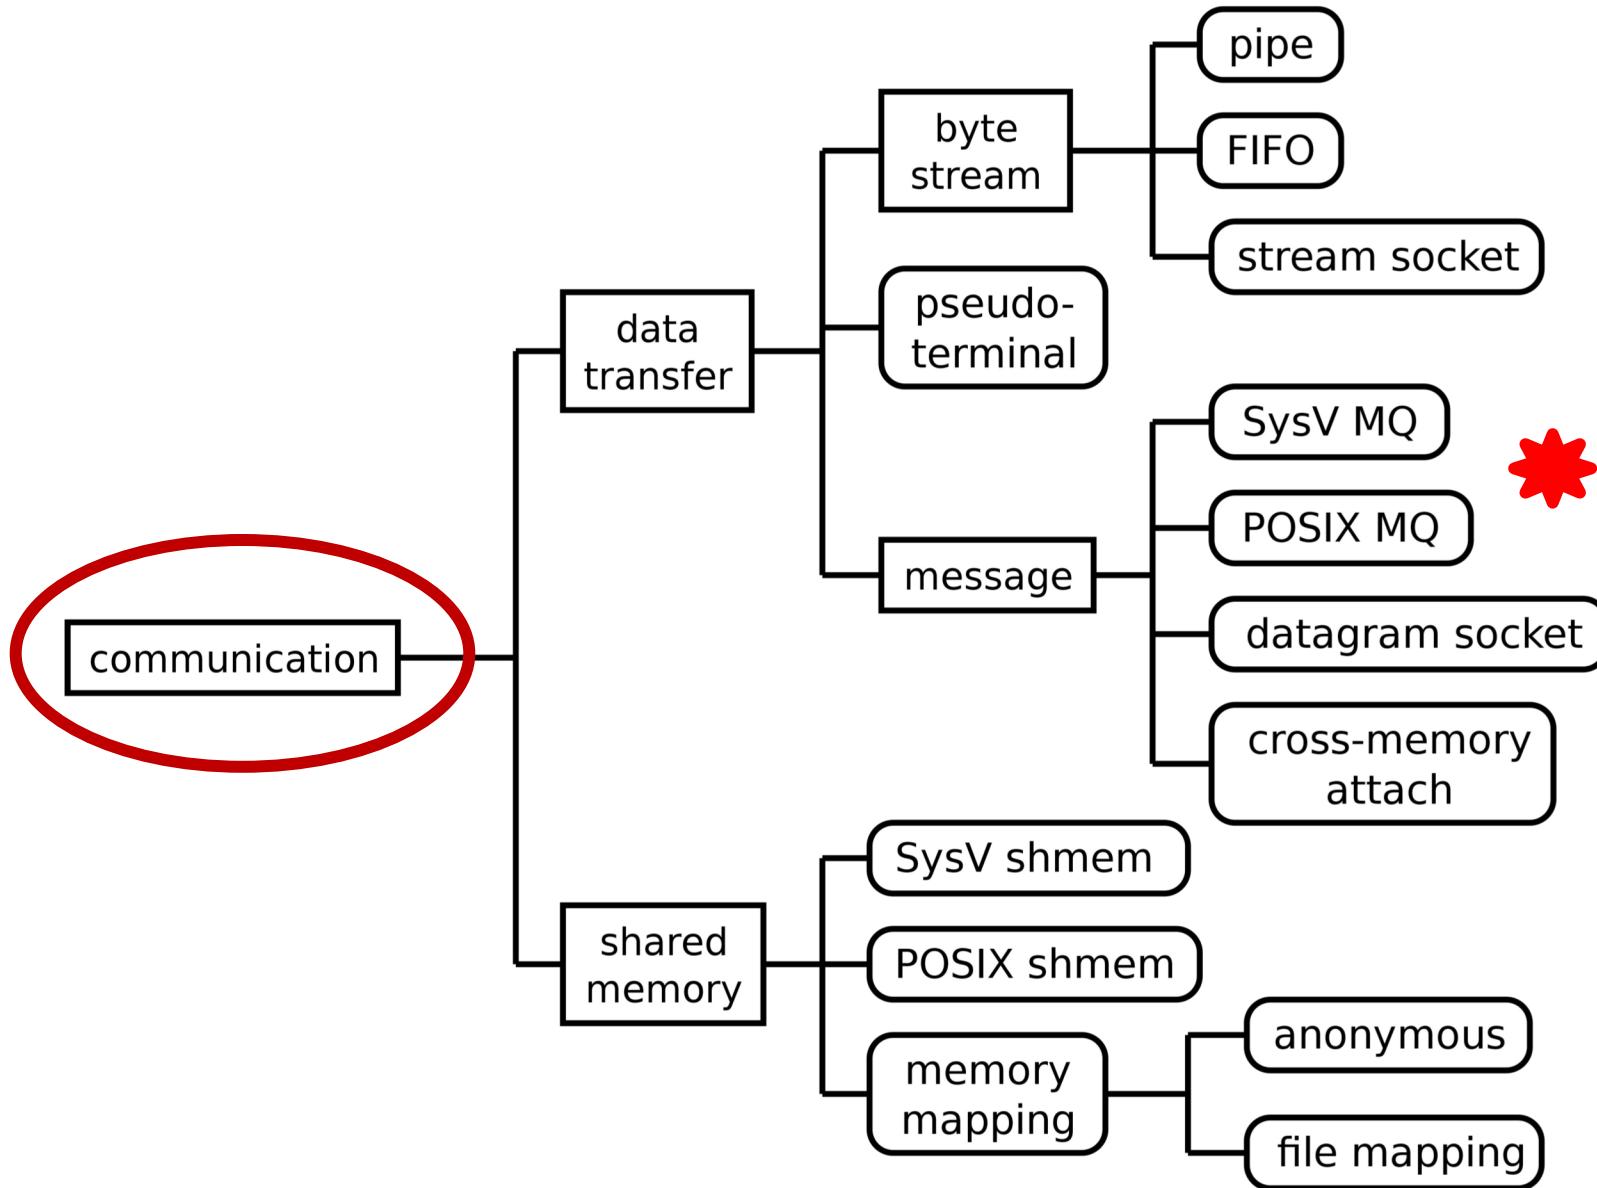
# Connecting two processes



# Basic IPC communication - Local

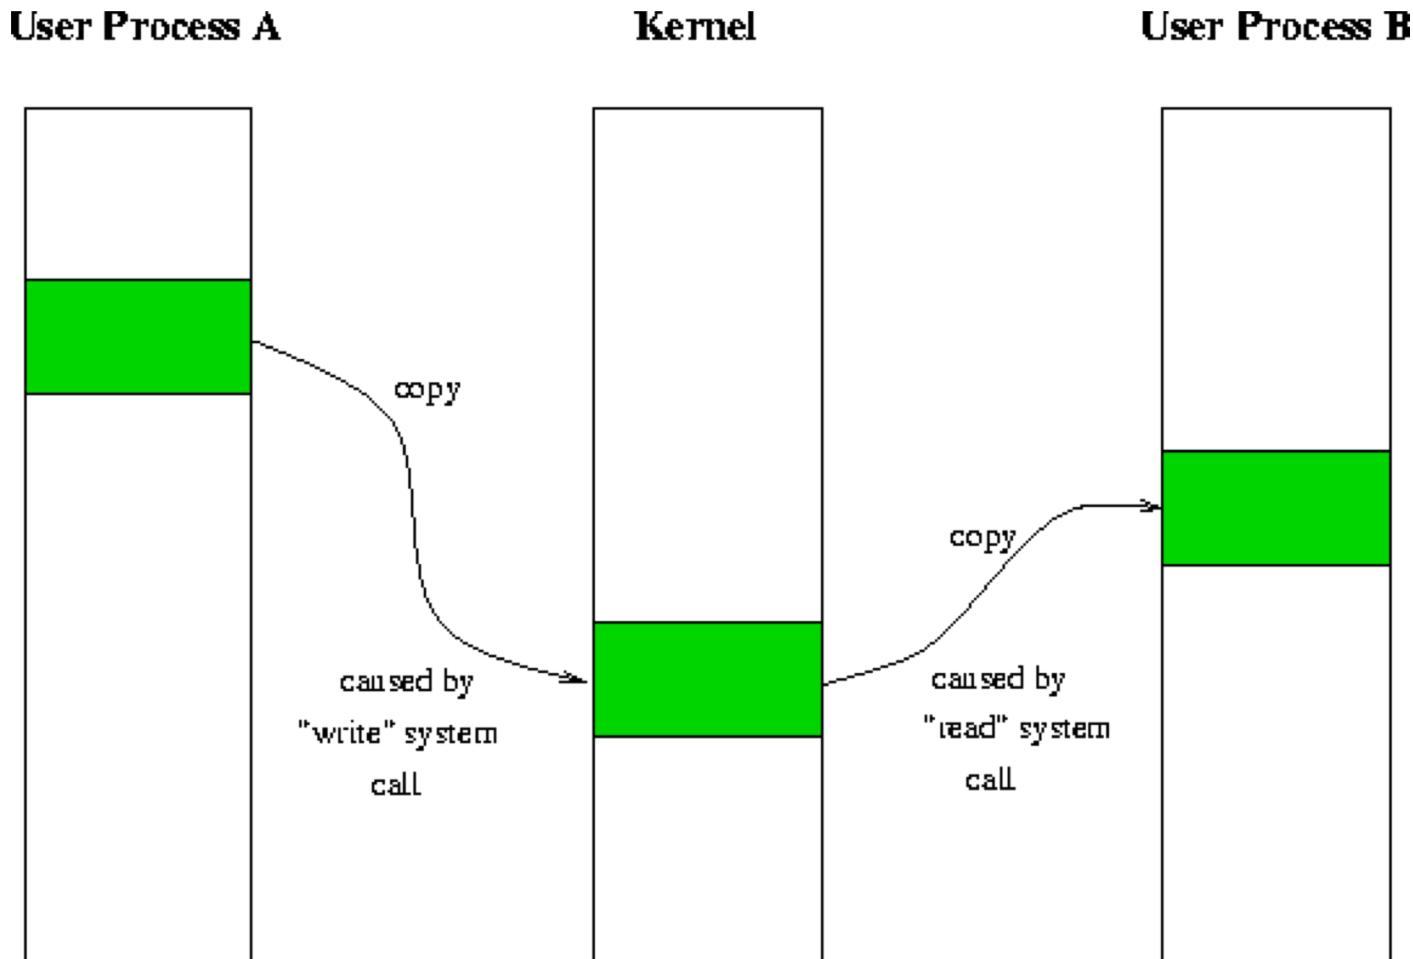
- Address space of different processes – independent
- The kernel has access to all memory
- Kernel copies the data first from the user address space into some internal kernel buffers using write() system call
- Then, kernel copies data from kernel buffers to the address space of another process
  - Caused by the read() system call from the second process

# Communication



# Basic IPC communication - Local

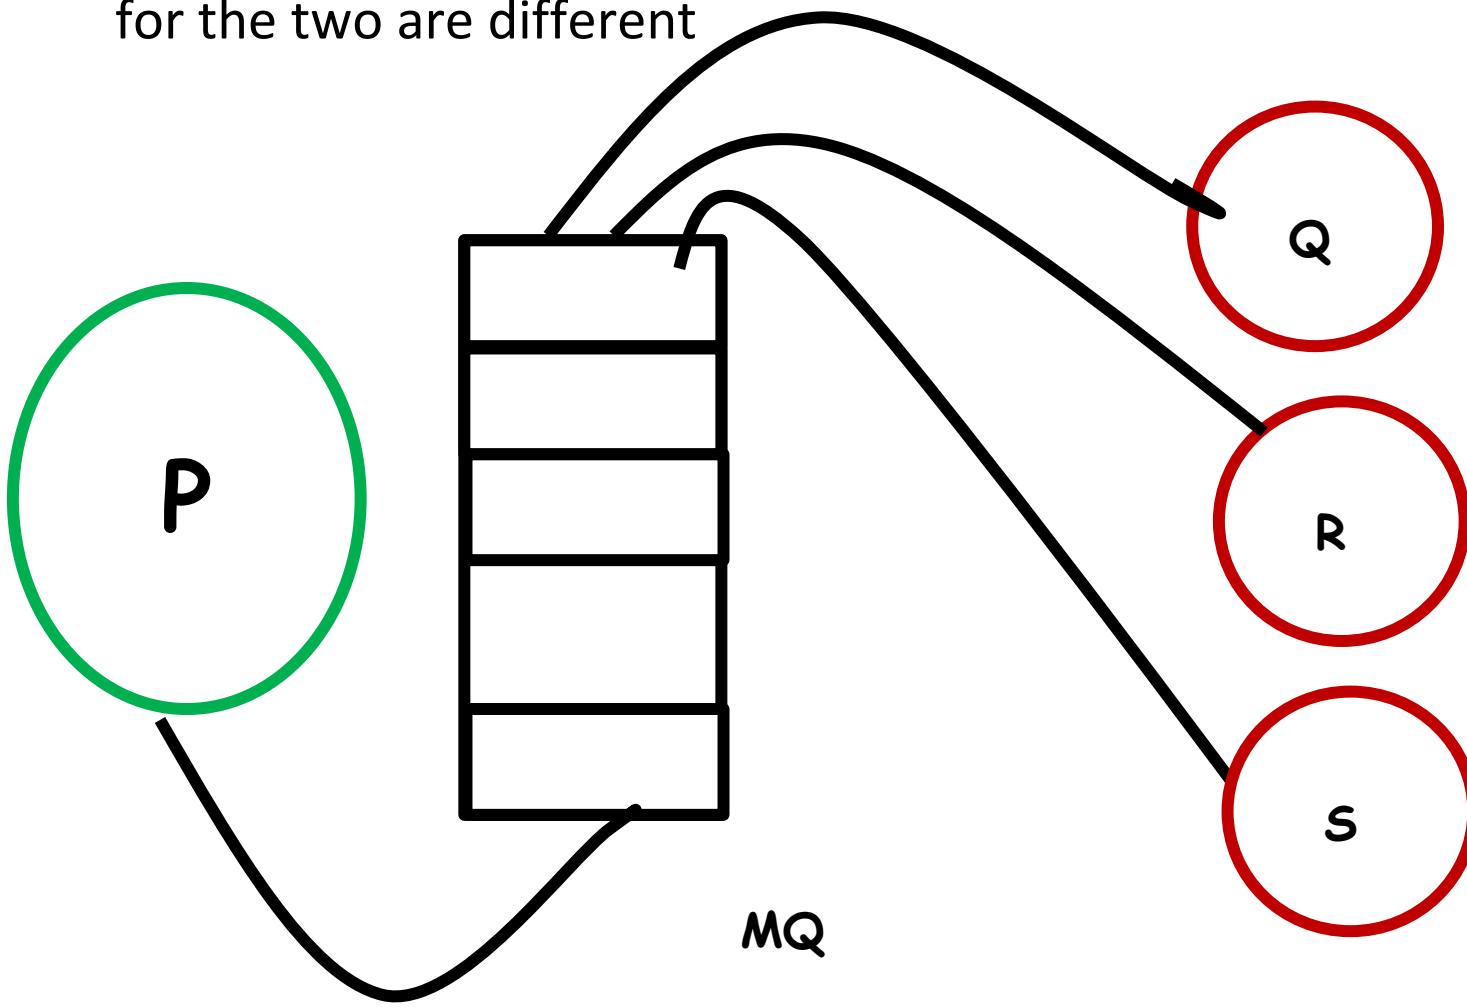
- Local (within a computer) interprocess communication uses memory for exchange of information



# **Linux / Unix Message Queues**

# Linux: Message Queue

- There are two varieties of message queues, System V message queues and POSIX message queues
  - Both provide almost the same functionality but system calls for the two are different



# Message Queue: 1 – Key generation

- To create a message queue, we need a System V IPC key
- We can create the key with the `ftok` function

```
#include <sys/types.h>
#include <sys/ipc.h>

key_t ftok (const char *pathname, int proj_id);
```

- The *pathname* must be an existing and accessible file
  - Any file name will do; any project id will do
  - However, these are known to other processes that want to communicate

# Message Queue: 2 – MSQ identifier

- Note that the message queue, semaphore and shared memory, has an associated system-wide identifier
- A process knowing this identifier, and having the relevant permission, can use the queue (or others)
- The `msgget` system call gets the message queue identifier for the given *key*

```
int msgget (key_t key, int msg_flags);
```

- If the `IPC_CREAT` flag is set in *msg\_flags*, the queue is created
  - Queue permissions are expressed as nine bits comprising of read, write and execute for owner, group and others

# Message Queue: 3 – Control Part

- With `msgctl`, we can do control operations on a message queue identified by `msqid`

```
int msgctl (int msqid, int cmd, struct msqid_ds *buf);
```

- The `cmd` parameter identifies the operation to be done
  - Removal of the queue
  - Change properties of the queue etc

# Message Queue: 4 – Message Send

- The `msgsnd` system call is used to sending messages to a System V message queue

```
int msgsnd (int msqid, const void *msgp, size_t msgsz, int msgflg);
```

- `*msg` : pointer to the message
- In most cases, the value of `msgflg` would be zero
  - Special cases are specified

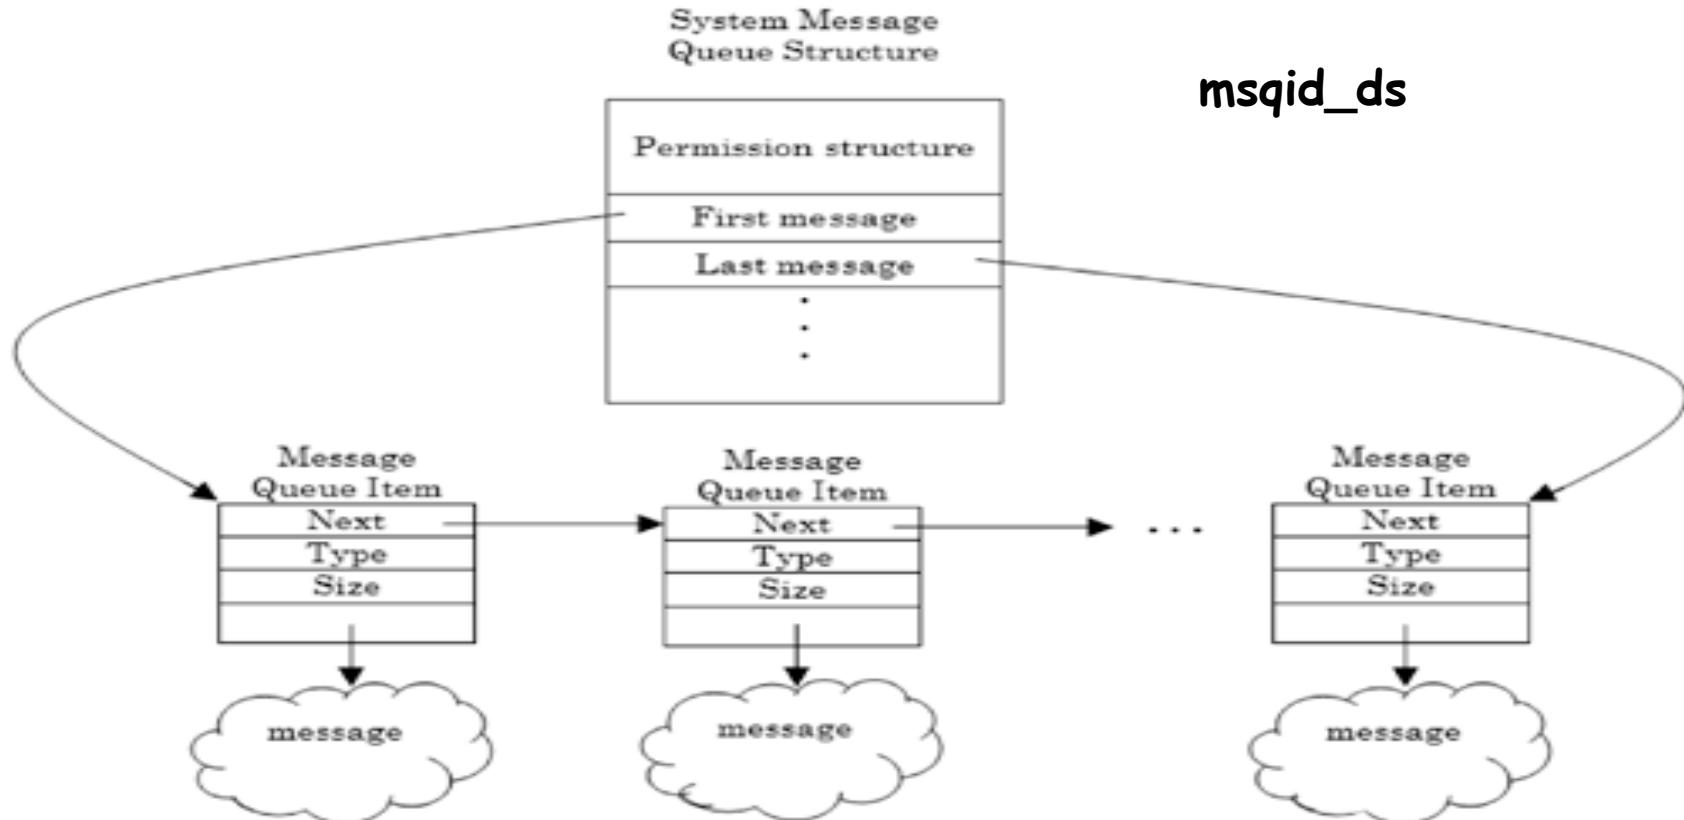
# Message Queue: 5 – Message Receive

- The `msgrcv` system call is for receiving messages from a message queue identified by `msqid`
- `msgp` points to the buffer for incoming message and `msgsz` specifies the maximum space available for `message_text` member of the message

```
ssize_t msgrcv (int msqid, void *msgp, size_t msgsz, long msgtyp, int msgflg);
```

# Message Queue – Data structure

- A new queue is created or an existing queue opened by **msgget()**
- All processes can exchange information through access to a common system message queue
- Each message is given an identification or type so that processes can select the appropriate

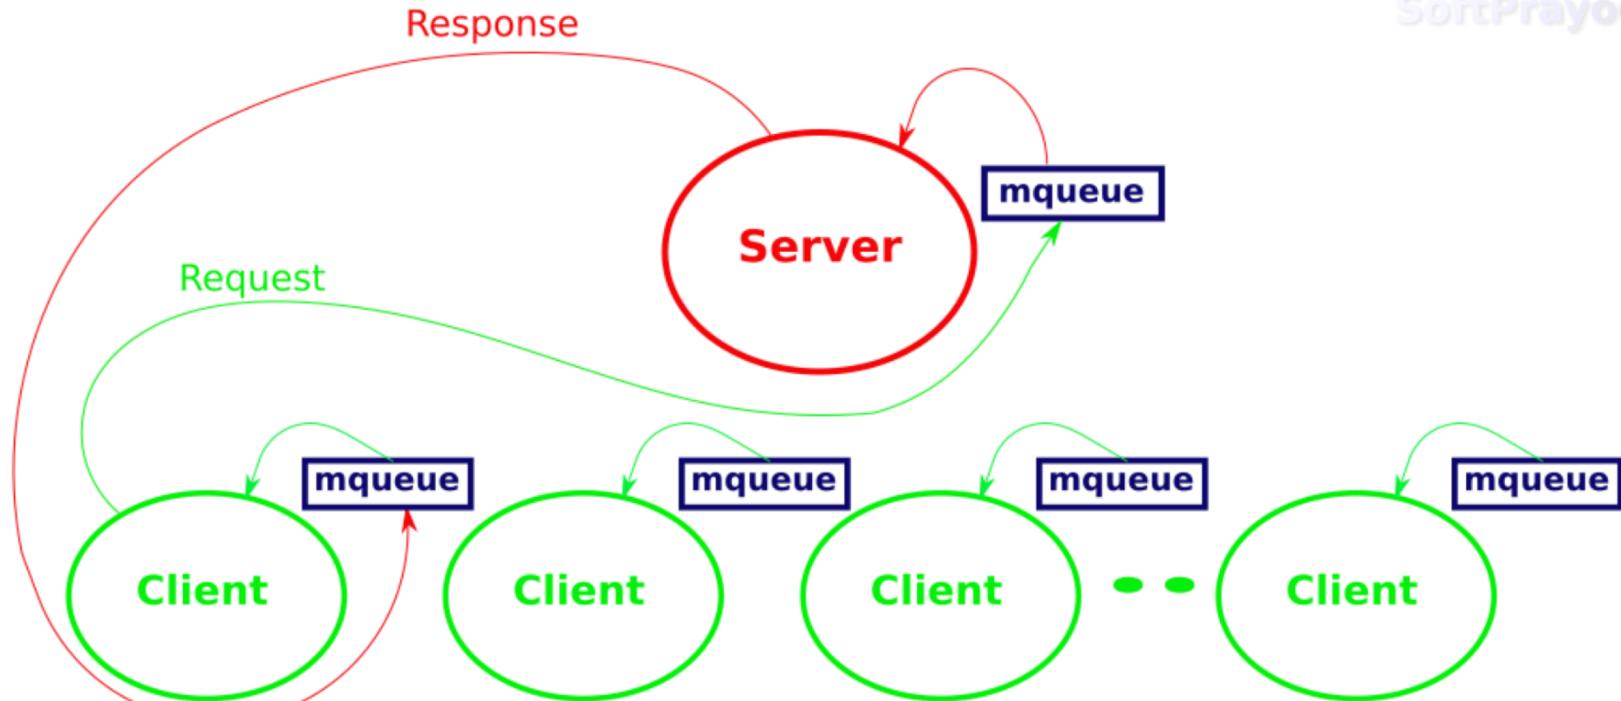


# Message Queue: Example

- The server listens for requests from clients
- The server has a message queue
  - Note that the parameters to the `ftok` function are known to clients (co-operating processes)
- The clients use these parameters to get the server queue identifier for communicating with the server
- Server queue allows reading operation for the server and writing operations by client
- The clients create their queue with the key `IPC_PRIVATE` and embed their queue identifier in the request message to the server
- So now both clients and server know each others queue identifiers

# Message Queue Example

SoftPrayog

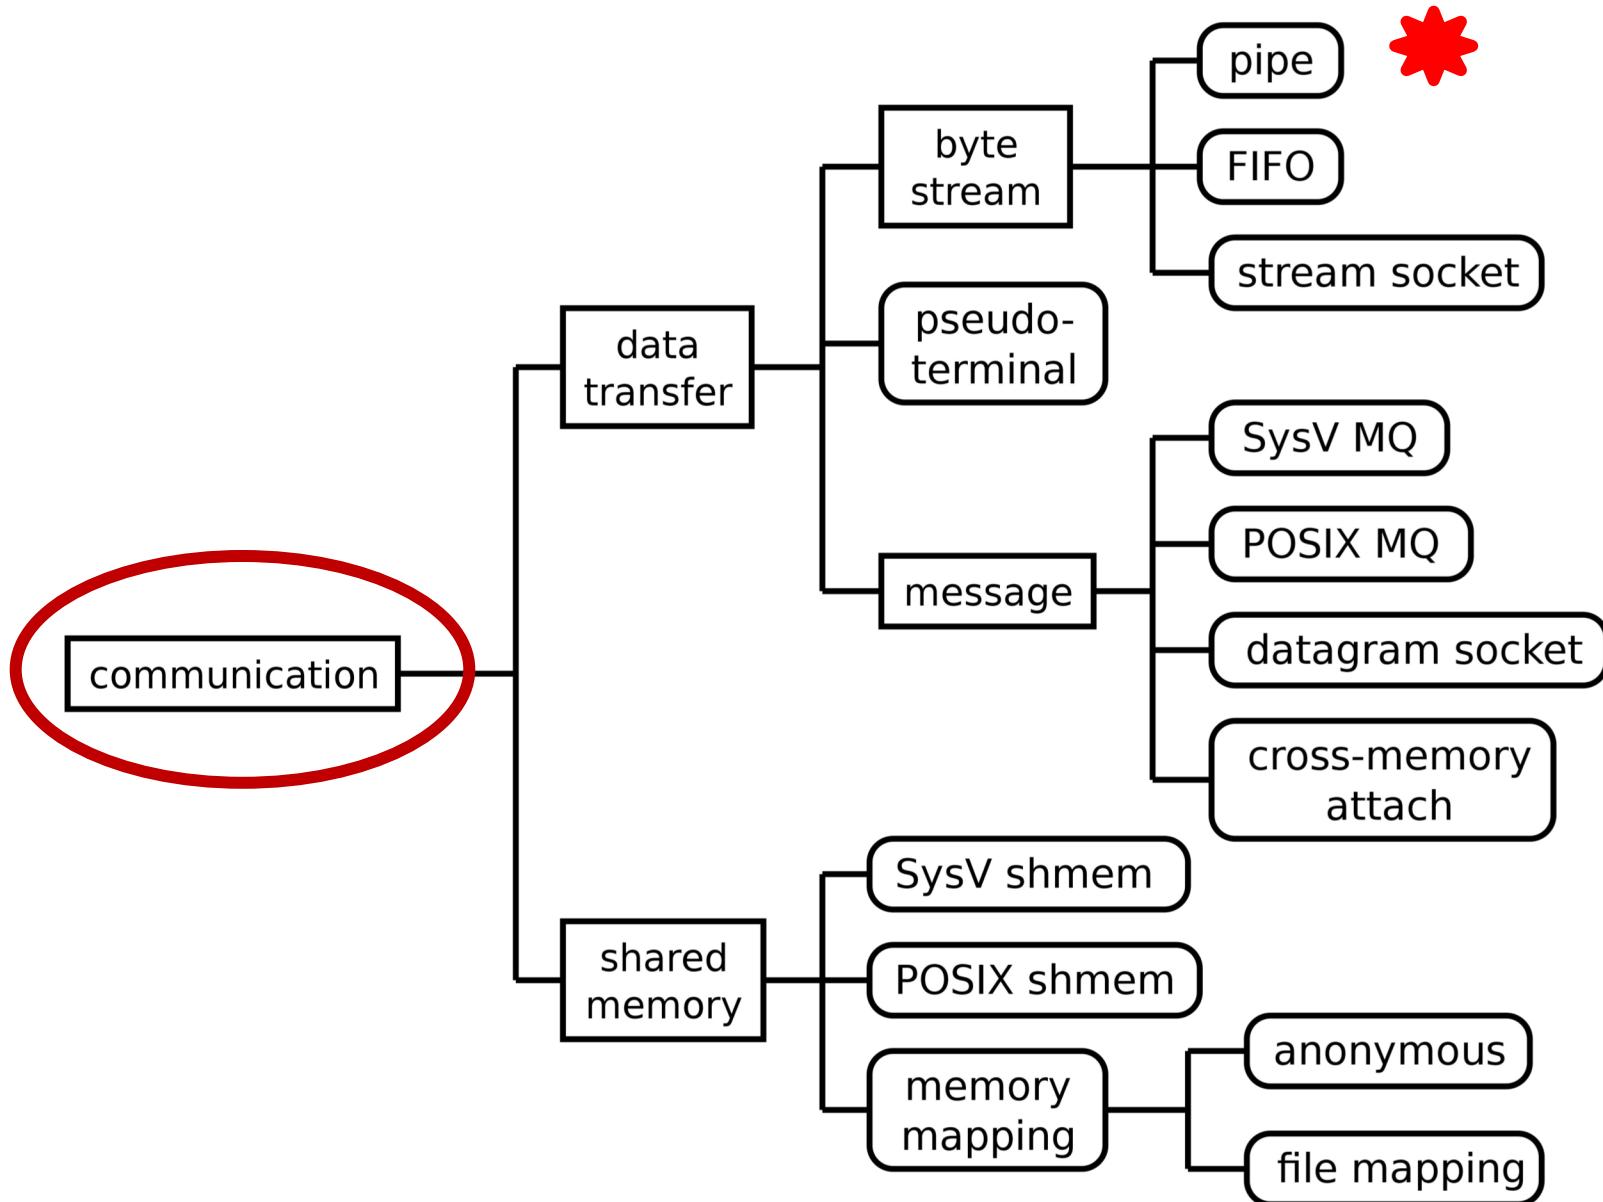


Interprocess Communication between client and server using Message Queues

# Message Queue Operations - Posix

- `mq_open()`: open/create MQ, set attributes
- `mq_close()`: close MQ
- `mq_send()`: send message
  - blocks if queue is full
- `mq_receive()`: receive message
  - blocks if no messages in queue
- `mq_unlink()`: remove MQ pathname
- `mq_setattr()`, `mq_getattr()`: set/get MQ attributes
- `mq_notify()`: request notification of msg arrival

# Communication



# Pipes

# Linux Pipes

- A *pipe* is a one-way flow of data between processes
- A process writes data to pipe
- Data is routed to another process by kernel
- The recipient process can read the data
- Take everything from standard out of program1 and pass it to standard input of program2
  - `program1 | program2`
- In Unix/Linux shells, pipes can be created by means of the `|` operator

# Linux Pipes

- Classic IPC method under UNIX:

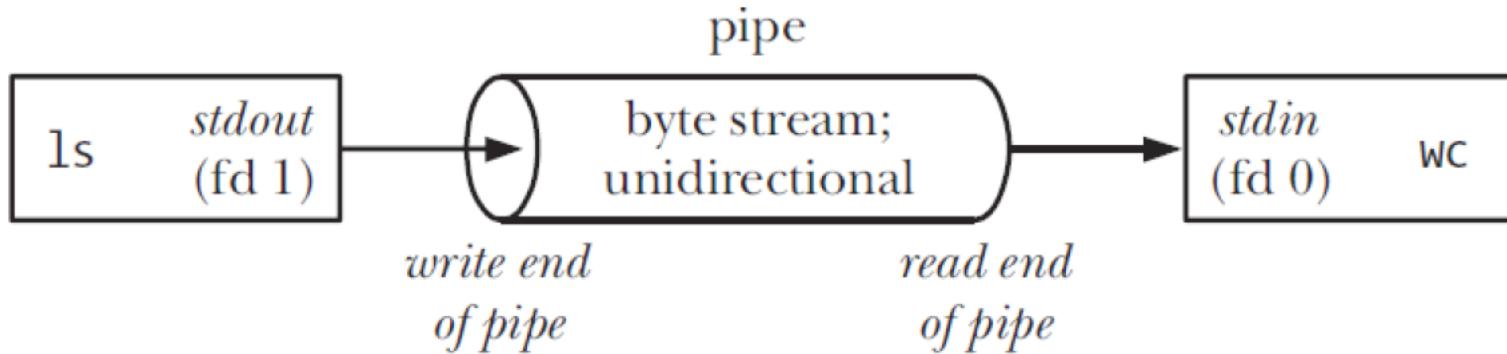
```
$ ls | wc -l
```

- shell runs two processes `ls` and `wc` which are linked via a pipe
- the first process (`ls`) writes data (e.g., using `write`) to the pipe and the second (`wc`) reads data (e.g., using `read`) from the pipe

- It is similar to

```
$ ls > temp
```

```
$ wc < temp
```



# Linux Pipe Implementation (self reading)

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the *write-end* of the pipe)
- Consumer reads from the other end (the *read-end* of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes
- Pipes may be considered as open files that have no corresponding image in the mounted filesystems
- Produces / consumes data through file syscalls `write()` and `read()`
- Data stored in kernel via `pipefs` filesystem

# Linux Pipe Implementation

- A process creates a new pipe by means of the `pipe( )` system call
  - Returns a pair of **file descriptors**
  - The process may then pass these descriptors to its descendants through `fork()`
  - The processes can read from the pipe by using the `read()` system call with the **first file descriptor**
  - The Processes can write into the pipe by using the `write()` system call with the **second file descriptor**
  - Each process must close one file descriptor before using the other

# Aside: File Descriptors

- We will study these— when we study file systems
- A Unix file is a sequence of  $m$  bytes
- All I/O devices are represented as files:
- Even the kernel is represented as a file:
  - `/dev/kmem` (kernel memory image)
  - `/proc` (kernel data structures)
- Basic Unix I/O operations (system calls):
  - `open()` and `close()`
  - `lseek()`
  - `read()` and `write()`

## Aside: File Descriptors

- Three file descriptors are already open when the process begins
  - File descriptor 0 is the *standard input*,
  - File descriptor 1 is the *standard output*
  - File descriptor 2 is the *standard error output*

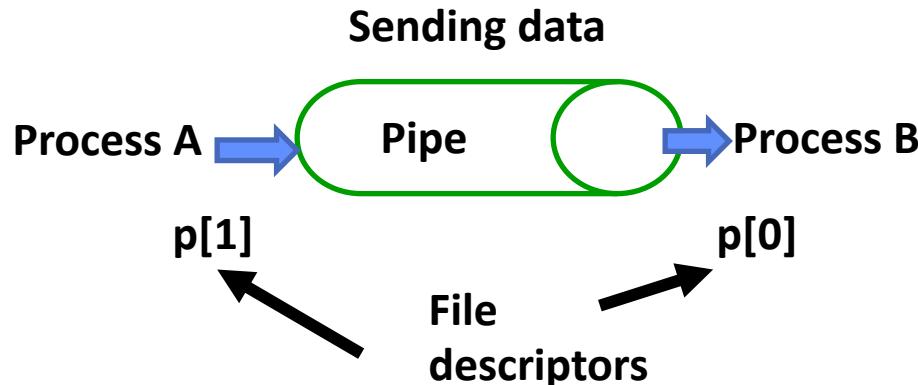
# Aside: File Descriptor

- Opening a file informs the kernel that you are getting ready to access that file

```
int fd; /* file descriptor */
if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {
    perror("open");
    exit(1);
}
```

- Returns **fd**: file descriptor (-1 on error) : A small integer
- Each process created by a Unix shell begins with three open files:
  - 0: standard input
  - 1: standard output
  - 2: standard error
- Must specify mode: O\_RDONLY, O\_WRONLY , O\_RDWR
- Other operations use file descriptor instead of explicit file name

# Pipe



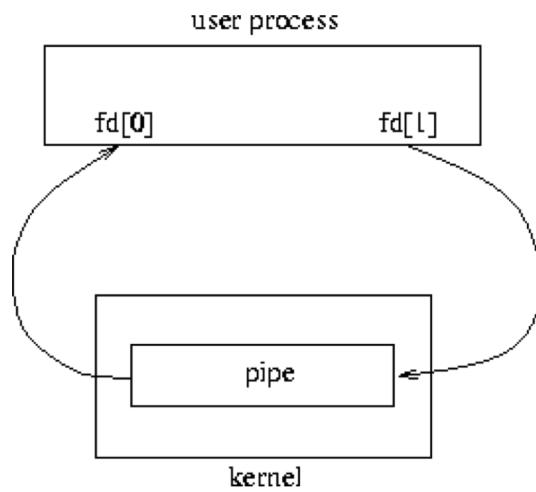
```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(void) {
    int pipefds[2];
    if(pipe(pipefds) == -1) {
        perror("pipe");
        exit(EXIT_FAILURE);
    }
    printf("Read File Descriptor Value: %d\n", pipefds[0]);
    printf("Write File Descriptor Value: %d\n", pipefds[1]);
    return EXIT_SUCCESS;
}
```

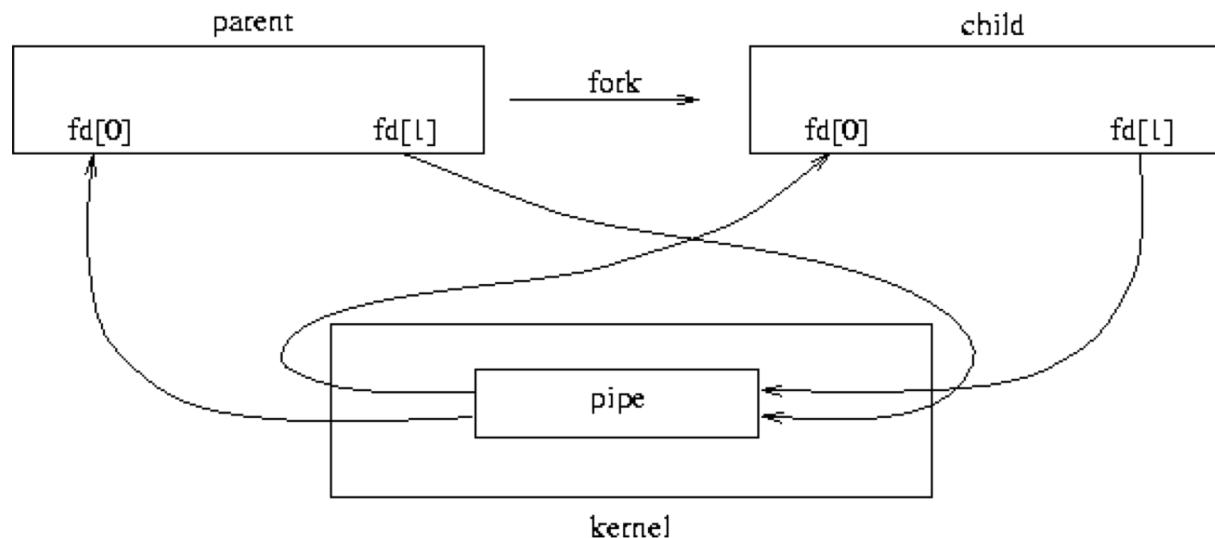
Pipe returns -1, if unsuccessful

Read File Descriptor Value: 3  
Write File Descriptor Value: 4

## The Result of the pipe() System Call

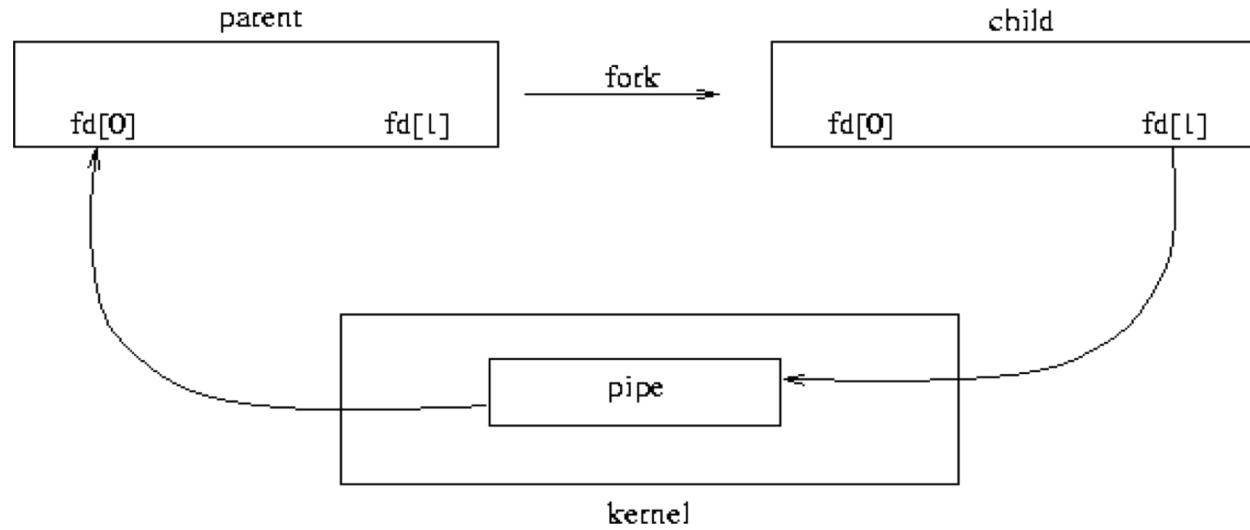


## A Pipe after a fork() System Call



## Resulting Pipe

### Pipe between Parent and Child Processes



# Can we implement ?

- ls | more with pipe()
- How can we attach stdout of ls with stdin of more ?

# Linux Pipe Implementation say `ls | more`

- Actions for `ls | more`
  - Shell invokes `pipe()` syscall
    - It returns file descriptors 3 (for reading) and 4 (for writing)
  - it invokes `fork()` system call **twice**
  - Releases file descriptor 3 and 4 (in Shell) with `close()`
- The first child process must execute `ls`
  - Invoke `dup2(4, 1)` to copy descriptor 4 to descriptor 1
    - Descriptor 1 now refers to pipe's write channel
    - Release descriptor 3 and 4 using `close()` syscall
  - Invoke `execve()` syscall to execute ls program
    - The program ls outputs to file with descriptor 1 i.e. writes to pipe

# Linux Pipe Implementation say `ls` | `more`

- The second child process must execute the `more` program
  - Invoke `dup2(3, 0)` to copy descriptor 3 to file descriptor 0;
    - Now onwards file descriptor 0 refers to pipe's read channel (from where reading can be done)
  - Releases file descriptor 3 and 4 using invoking `close()` system call, twice
  - Invokes `execve()` system call to execute `more`
    - By default, that program reads its input from the file that has file descriptor 0
- A pipe can be used by an arbitrary number of processes
  - However, this requires file locking for synchronization of reading/writing operations by multiple processes

## pipe() syscall

For writing

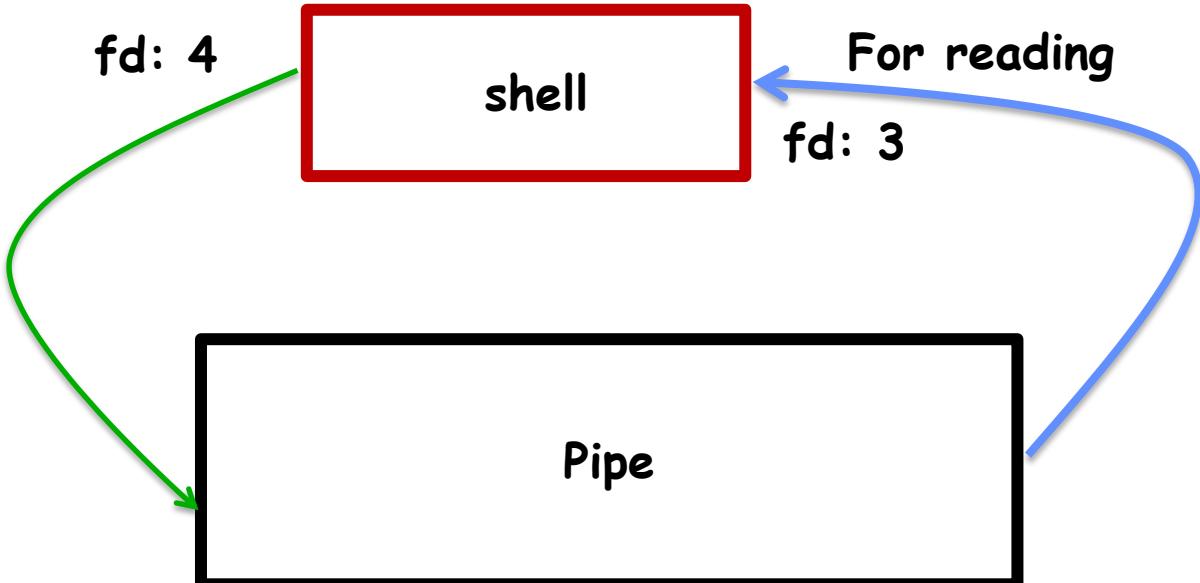
fd: 4

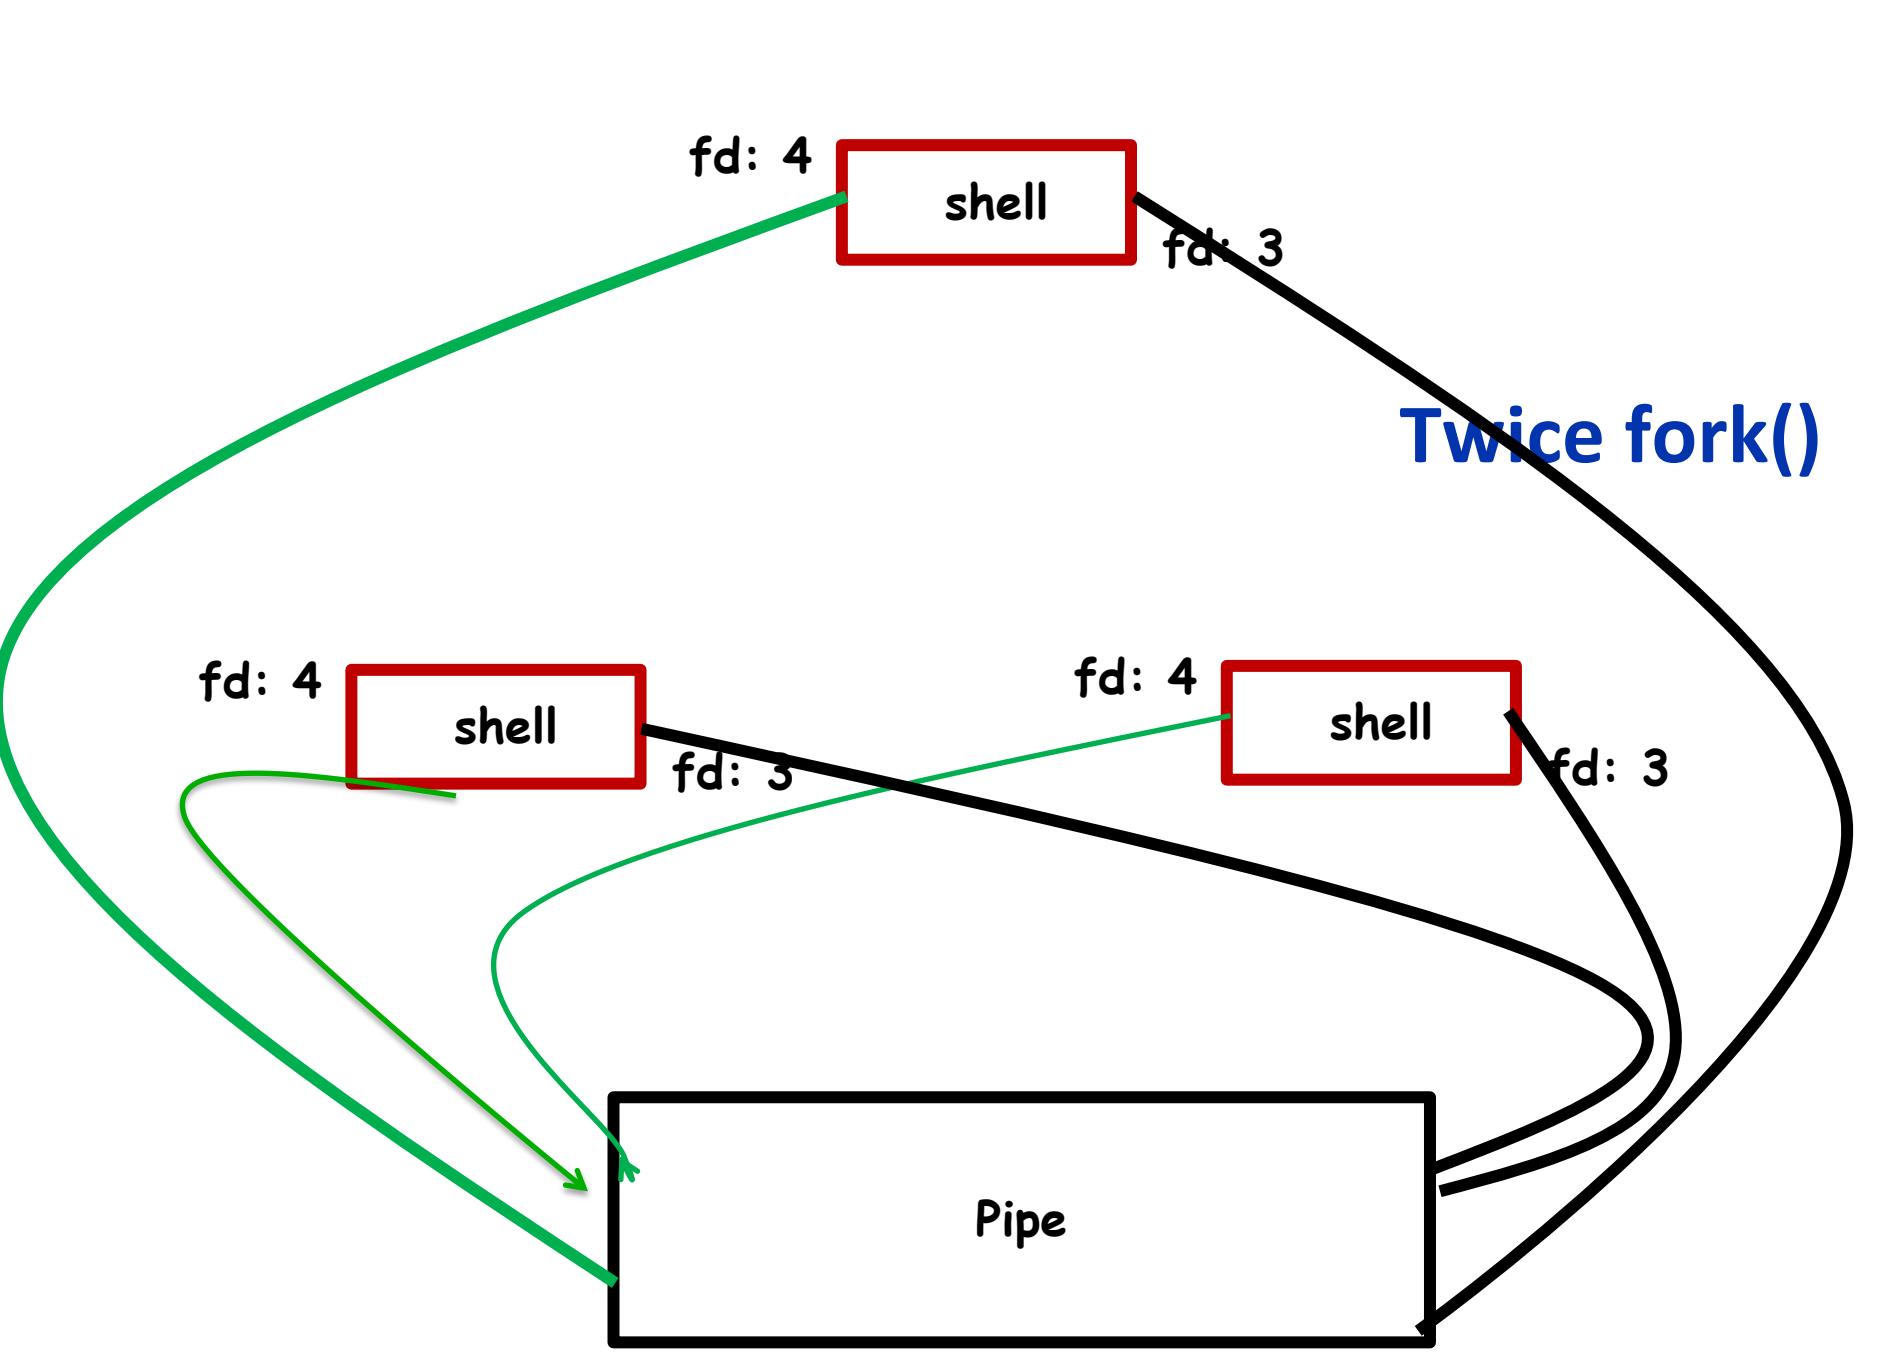
shell

For reading

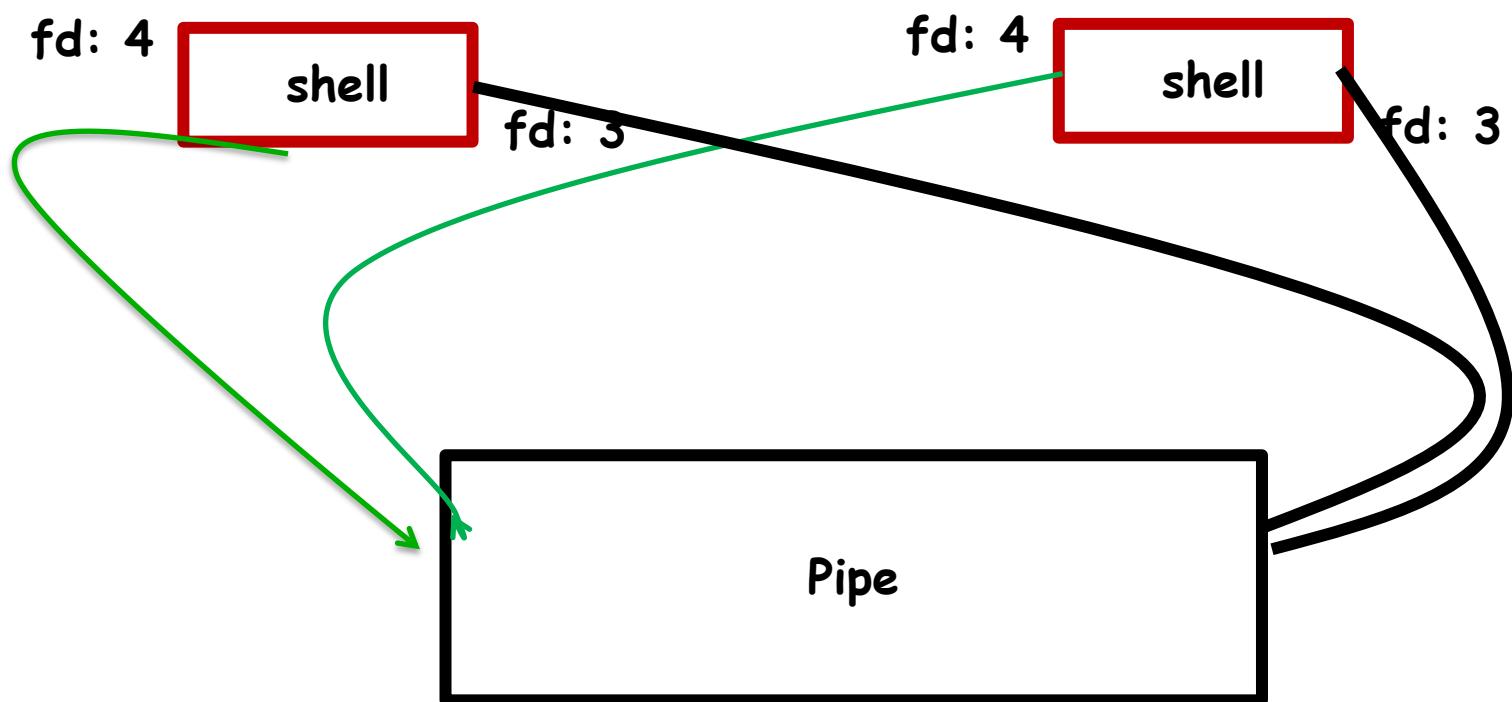
fd: 3

Pipe

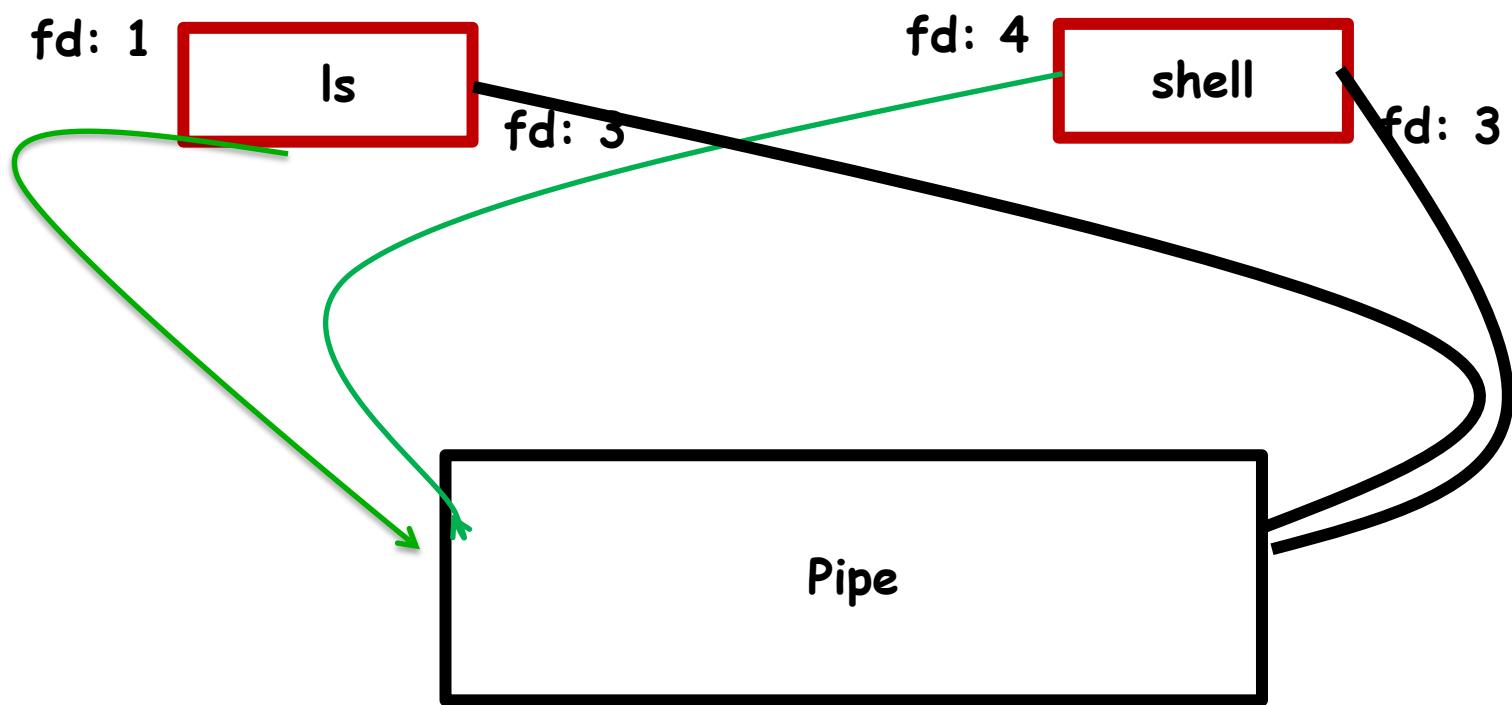




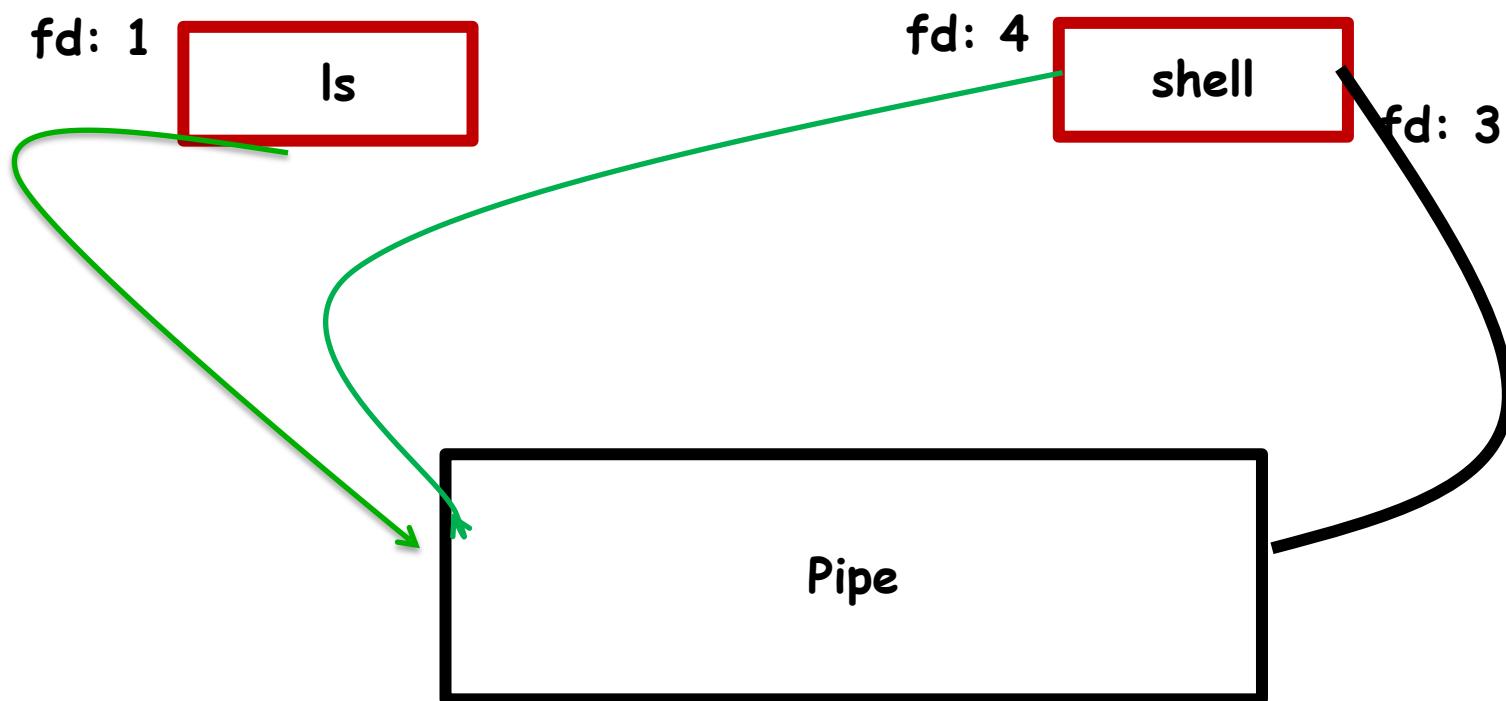
## Twice fork()



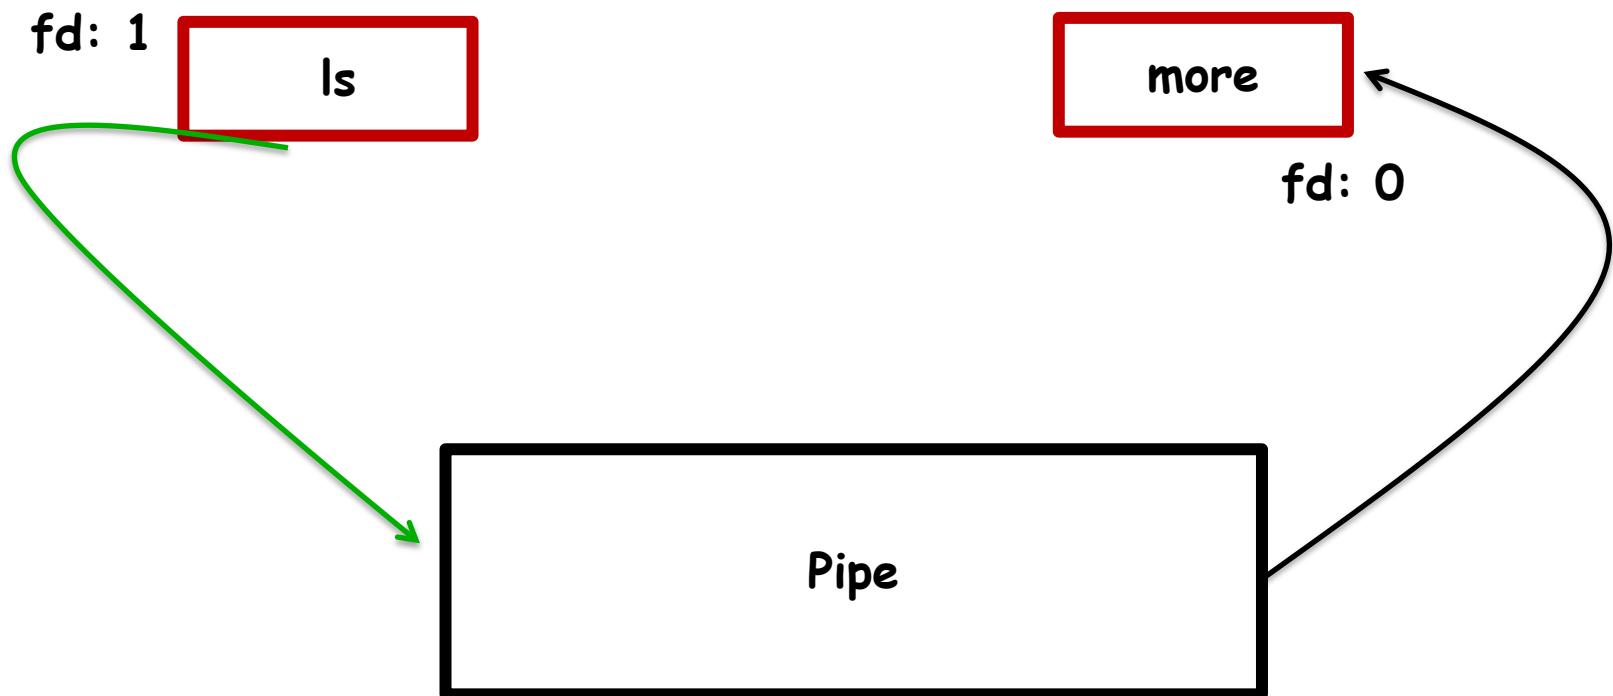
## Twice fork()



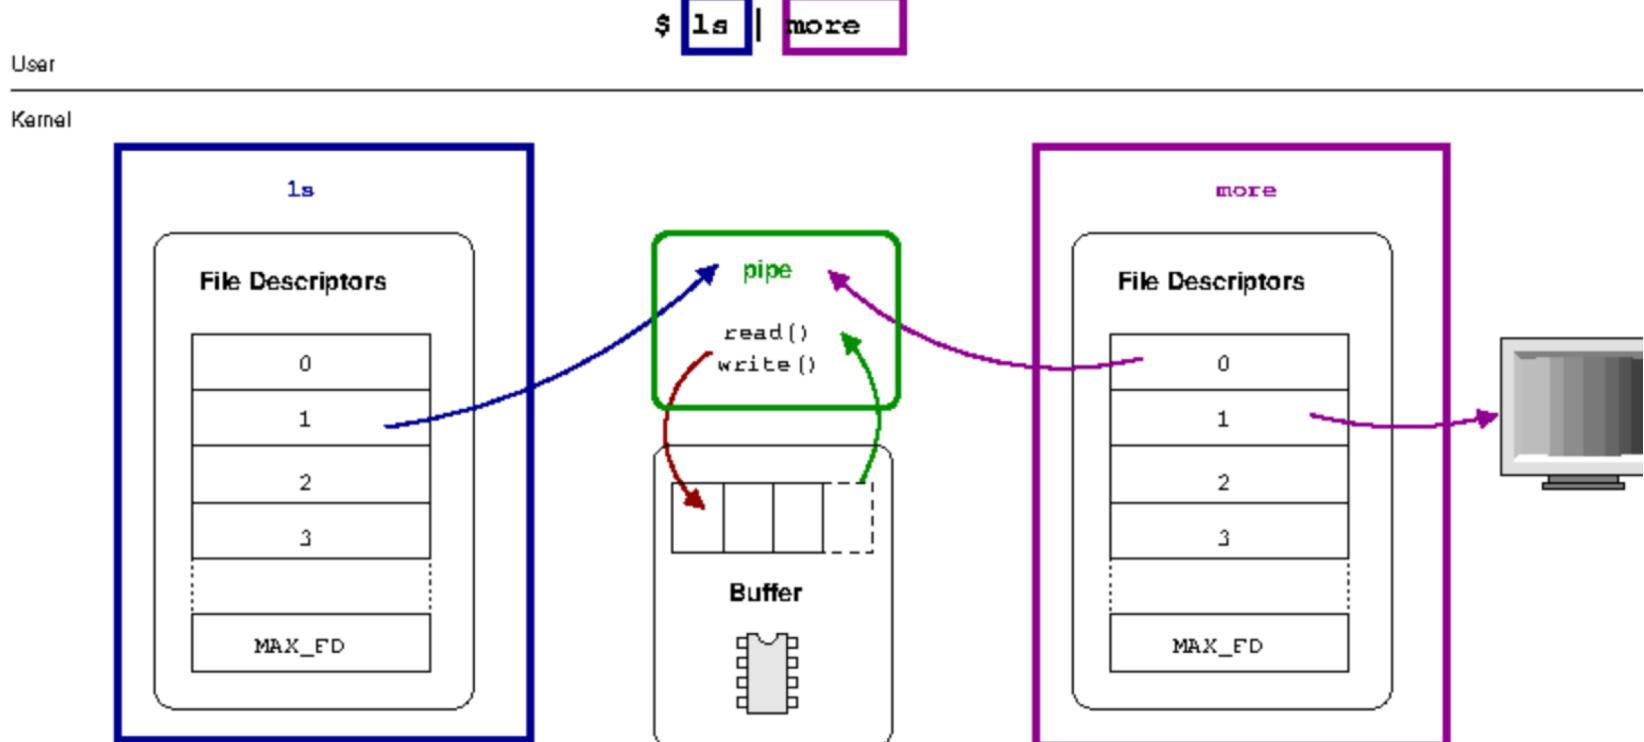
## Twice fork()



## Twice fork()



# Pipe Implementation



# Disadvantages of Pipe

- Pipe is created by a common ancestor
  - Impossible for two arbitrary processes to share the same pipe
  - Example: Database engine server and a client process
    - Clients can come and go
  - Note that Pipes can be used within a system
- Here comes
  - Named pipe or FIFO
- Ordinary pipes exist only while the processes are communicating with one another
  - Pipe is deleted when processes finished communication
- FIFO combines features of a regular file and a pipe
  - Unrelated processes can communicate with each other
  - Once opened FIFO behaves like a pipe

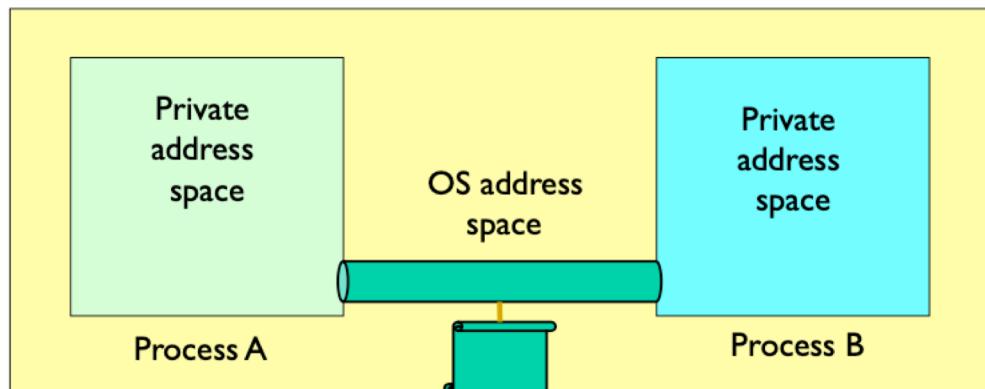
# **Named Pipes or FIFO (Self Reading)**

# Named File or FIFO (self reading)

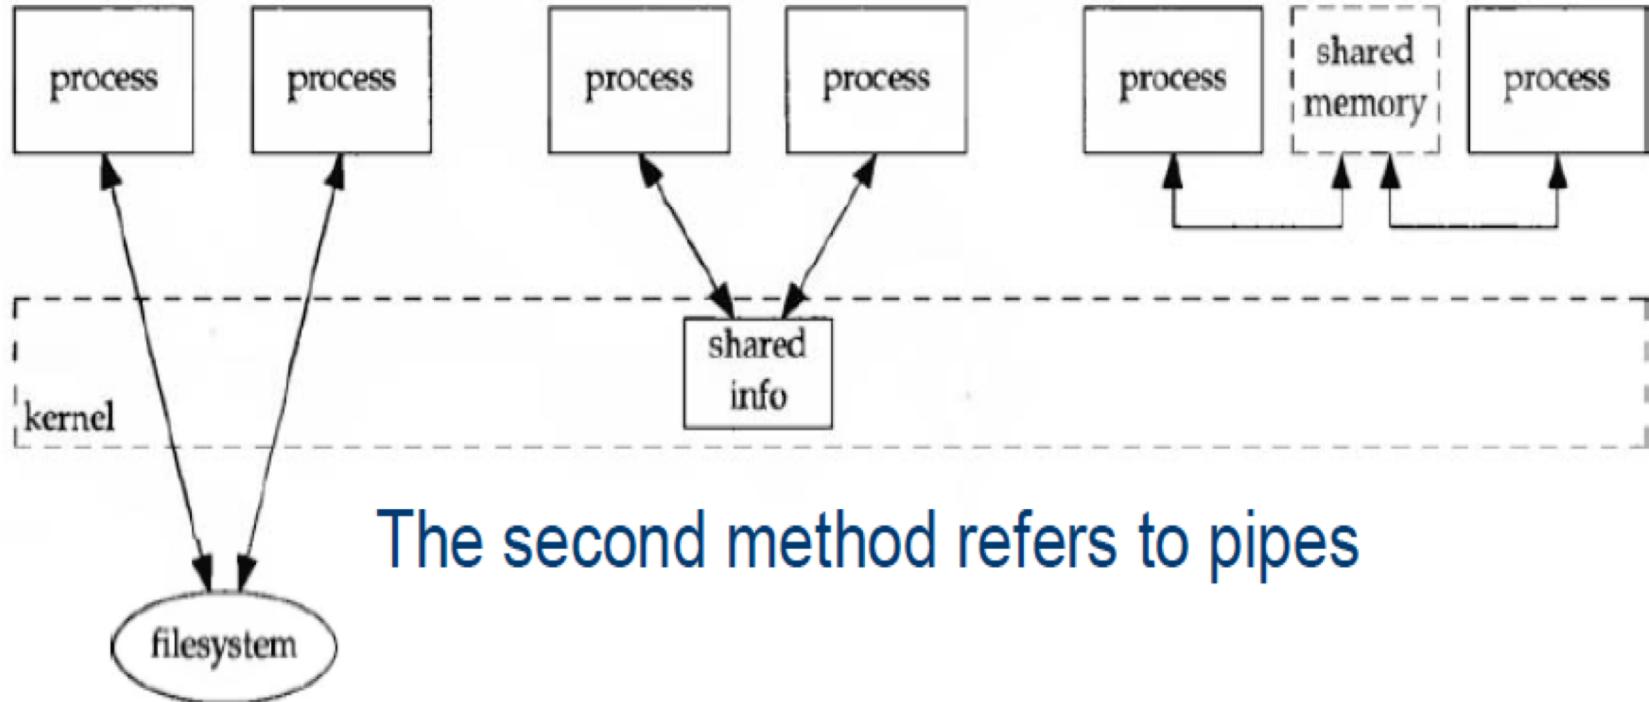
- FIFO is associated with kernel buffer that temporarily stores data exchanged by two or more processes
- Communication can be bidirectional
- No parent-child relationship required
- Once named pipe / FIFO is established several processes can use it for communication
- FIFOs are created by one process that calls mkfifo()
- FIFO is opened (twice) to get read and write file descriptors
- Now FIFO can be treated as a file

# Named Pipe or FIFO (self reading)

- Although FIFOs have a handle in the regular file system, they are not files
  - FIFOs are not backed by real file system
  - FIFOs must be opened for reading and writing before either may occur
  - Limited capacity
- Persistent – can be used multiple times



# Message Communication with Files

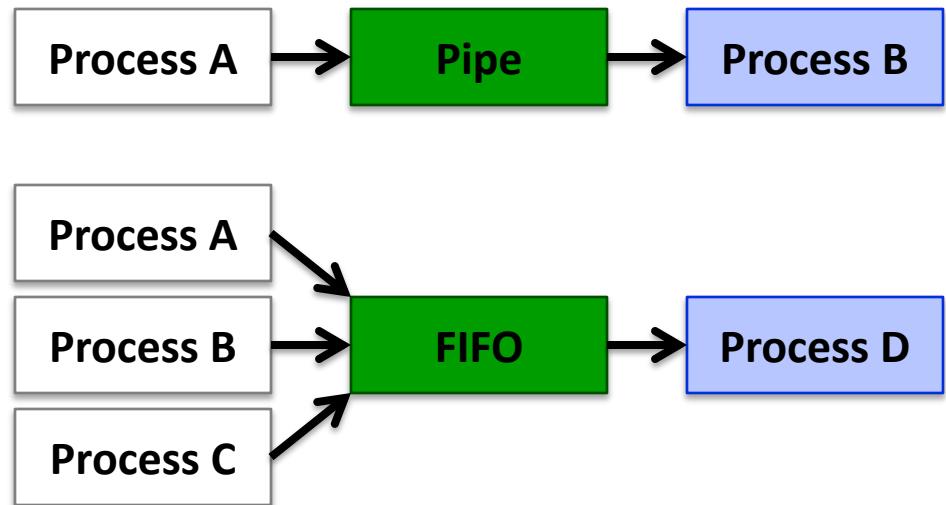


# Lecture Summary

Pipes are useful for implementing many design patterns and idioms:

Producer / Consumer

Client / Server



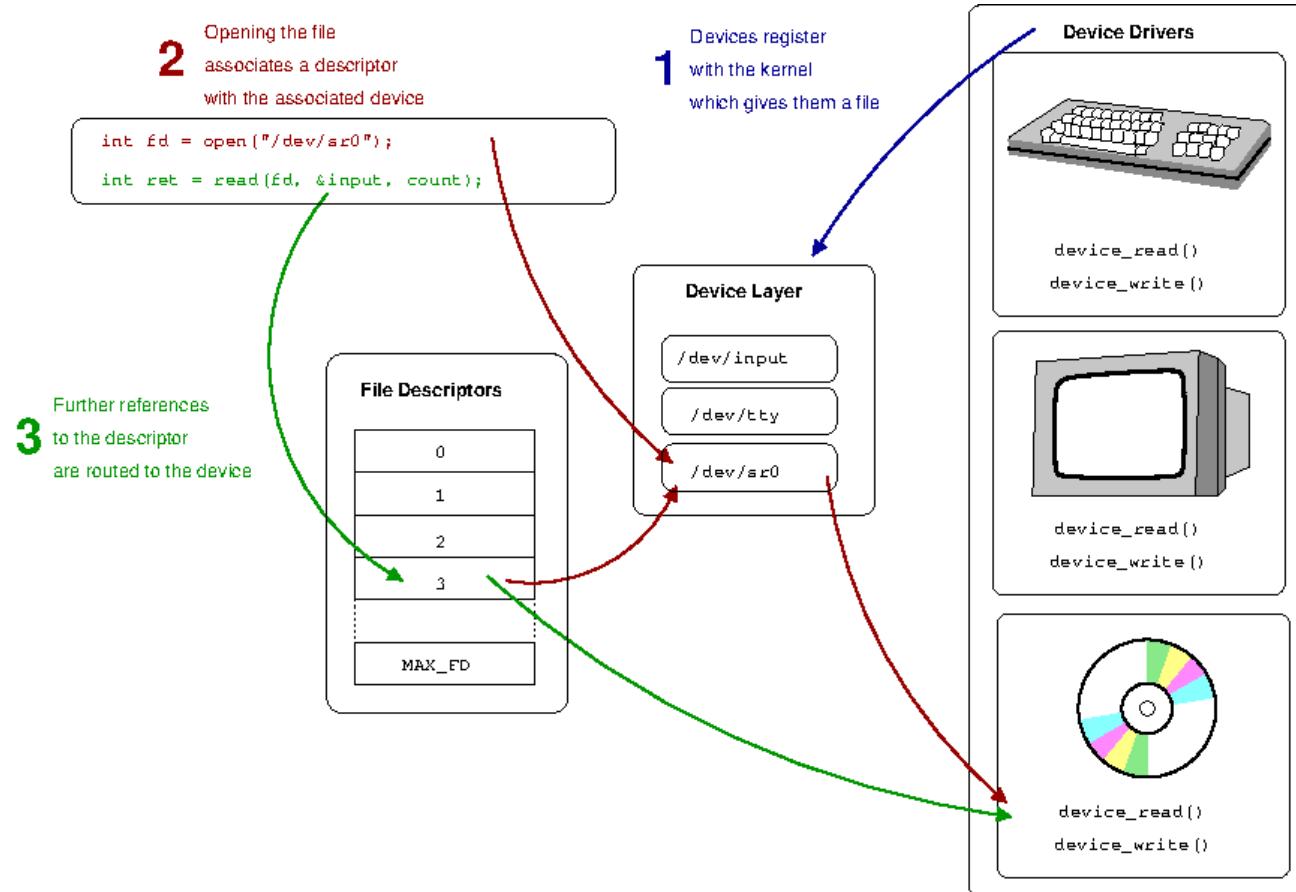
# Backup

# dup & dup2

```
#include <unistd.h>
int dup(int fildes);
```

- Returns a new file descriptor that is a copy of *fildes*
- File descriptor returned is first available file descriptor in file table.
- For example, to dup a read pipe end to stdin (0), close stdin, then immediately dup the pipe's read end.
- Close unused file descriptors; a process should have only one file descriptor open on a pipe end.

# File Descriptor (self reading)



# Linux: Message Queue

- A message queue is a linked list of messages stored within the kernel and identified by a message queue identifier
- Linux maintains a list of message queues, the **msgque** vector
  - Each element of **msgque** points to a **msqid\_ds** data structure that fully describes the message queue
- When message queues are created, a new **msqid\_ds** data structure is allocated from system memory and inserted into the **vector**
  - Each message
    - Type
    - Length
    - Data
  - message

# Posix Message Queue Attributes

```
struct mq_attr {  
    long mq_flags;      // MQ description flags  
                      // 0 or O_NONBLOCK  
                      // [mq_getattr(), mq_setattr()]  
    long mq_maxmsg;    // Max. # of msgs on queue  
                      // [mq_open(), mq_getattr()]  
    long mq_msgsize;   // Max. msg size (bytes)  
                      // [mq_open(), mq_getattr()]  
    long mq_curmsgs;   // # of msgs currently in queue  
                      // [mq_getattr()]  
};
```