

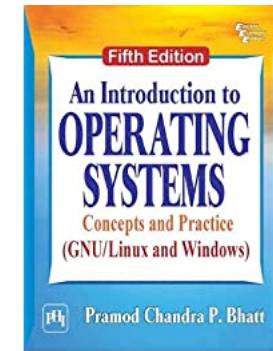
# **CS310 Operating Systems**

## **Lecture 18: Inter Process Communication – Shared Memory**

Ravi Mittal  
IIT Goa

# Reading

- Book: An introduction to Operating System: Concepts and Practices, Pramod Chandra Bhatt, 5<sup>th</sup> Edition,
  - Chapter 7, Section 7.3.4

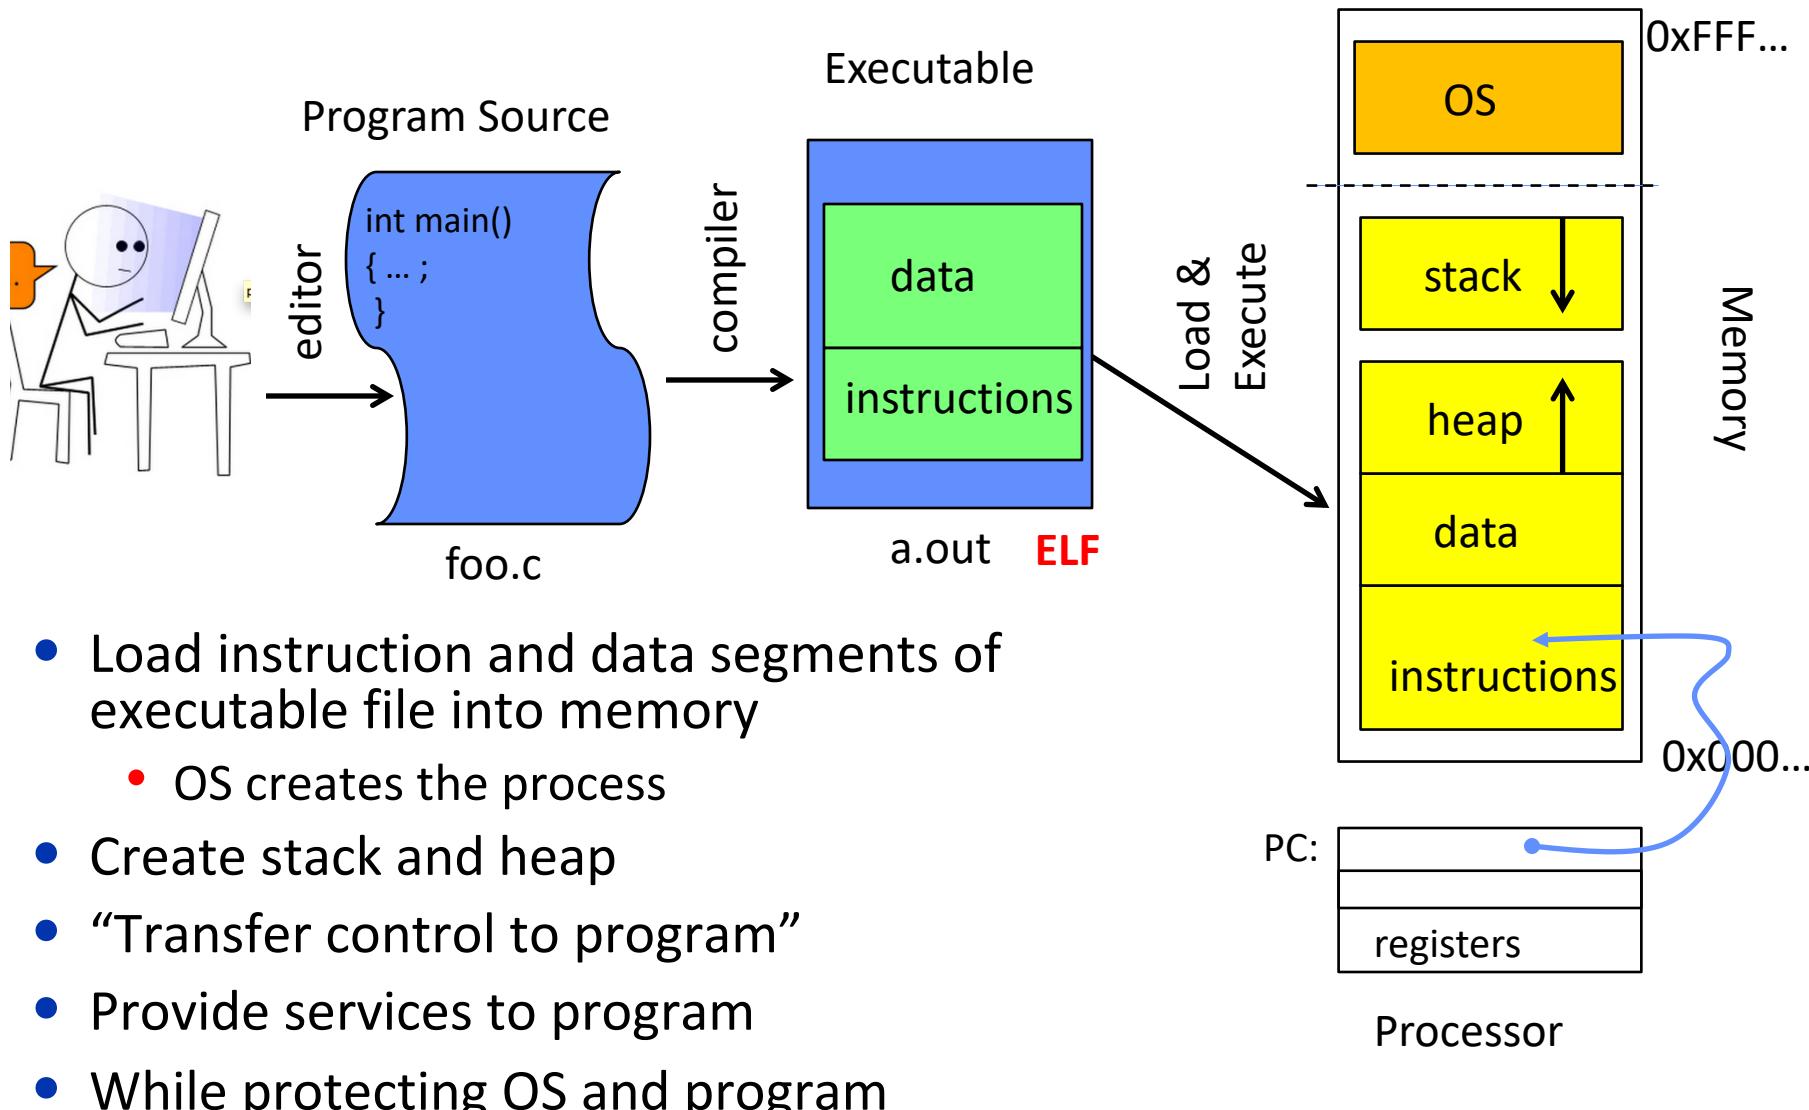


# Today we will study

- Where are we in CS310?
- Shared Memory Concept
- Garbage Collection

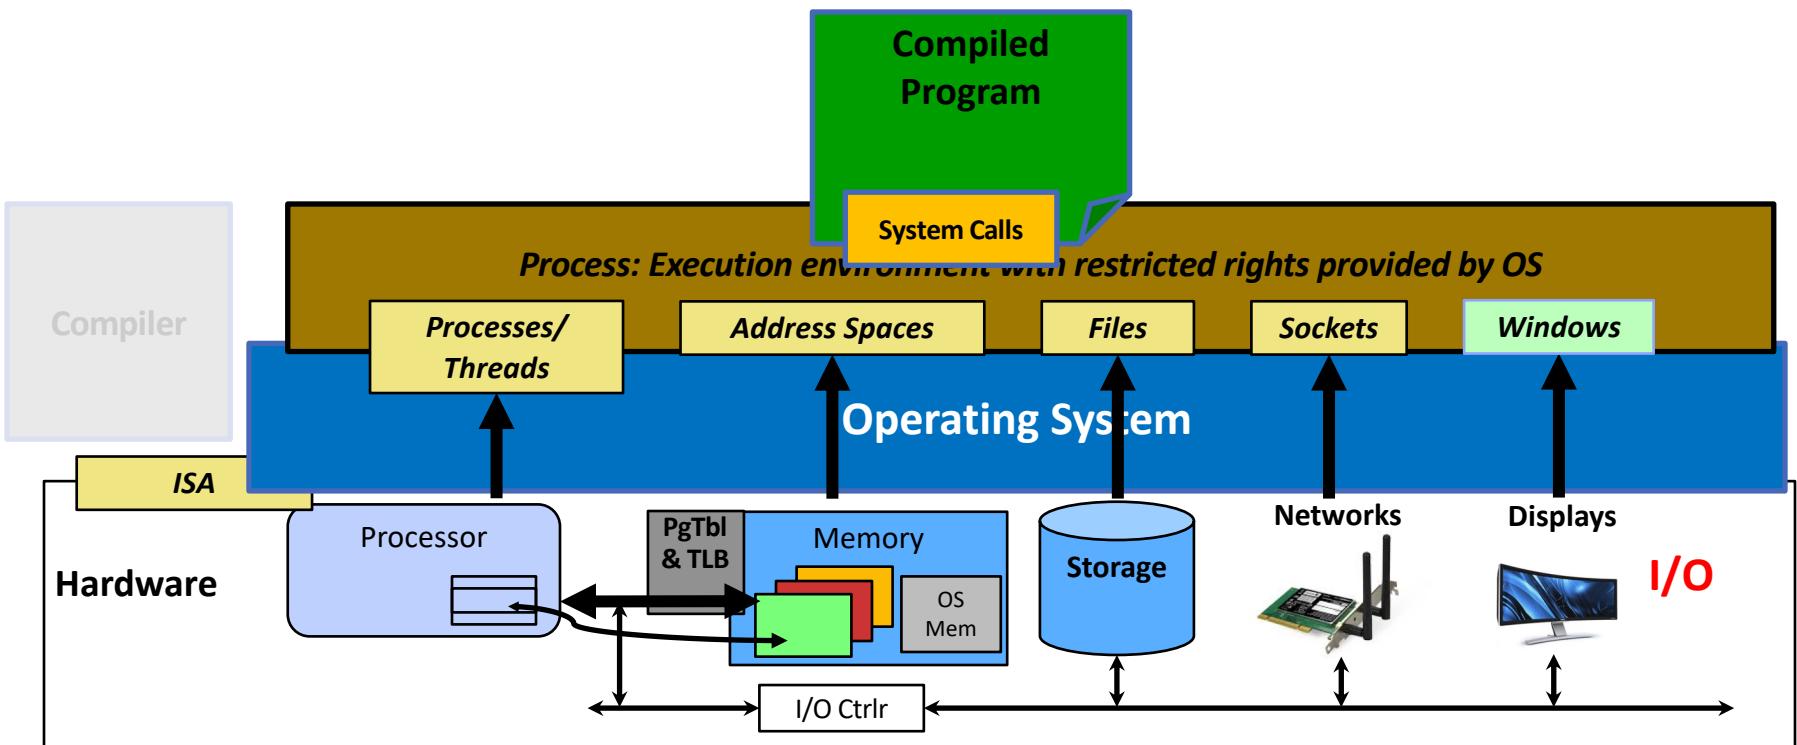
# **Where are we in CS310?**

# Recall: OS Bottom Line: Run Programs



**ELF: Executable and Linking Format**

# Recall: Operating System



# Topics

- OS Concepts: Top Level view
  - How does the OS work?
  - Multiprogramming Concept
- Process Concept, PCB, Context Switching
- How processes interact with OS for use of hardware – System Calls
- Threads
  - Concept
- Inter-process Communication
  - Message Passing and Shared Memory
- Concurrency Management
- File System, I/O
- Security
- Cloud Infrastructure

# Two important Requirements

- When processes/threads interact with each other, they must satisfy the following requirements:
  - Synchronization
    - Processes may need to be synchronized to enforce mutual exclusion
  - Communication
    - Cooperating processes may need to exchange information

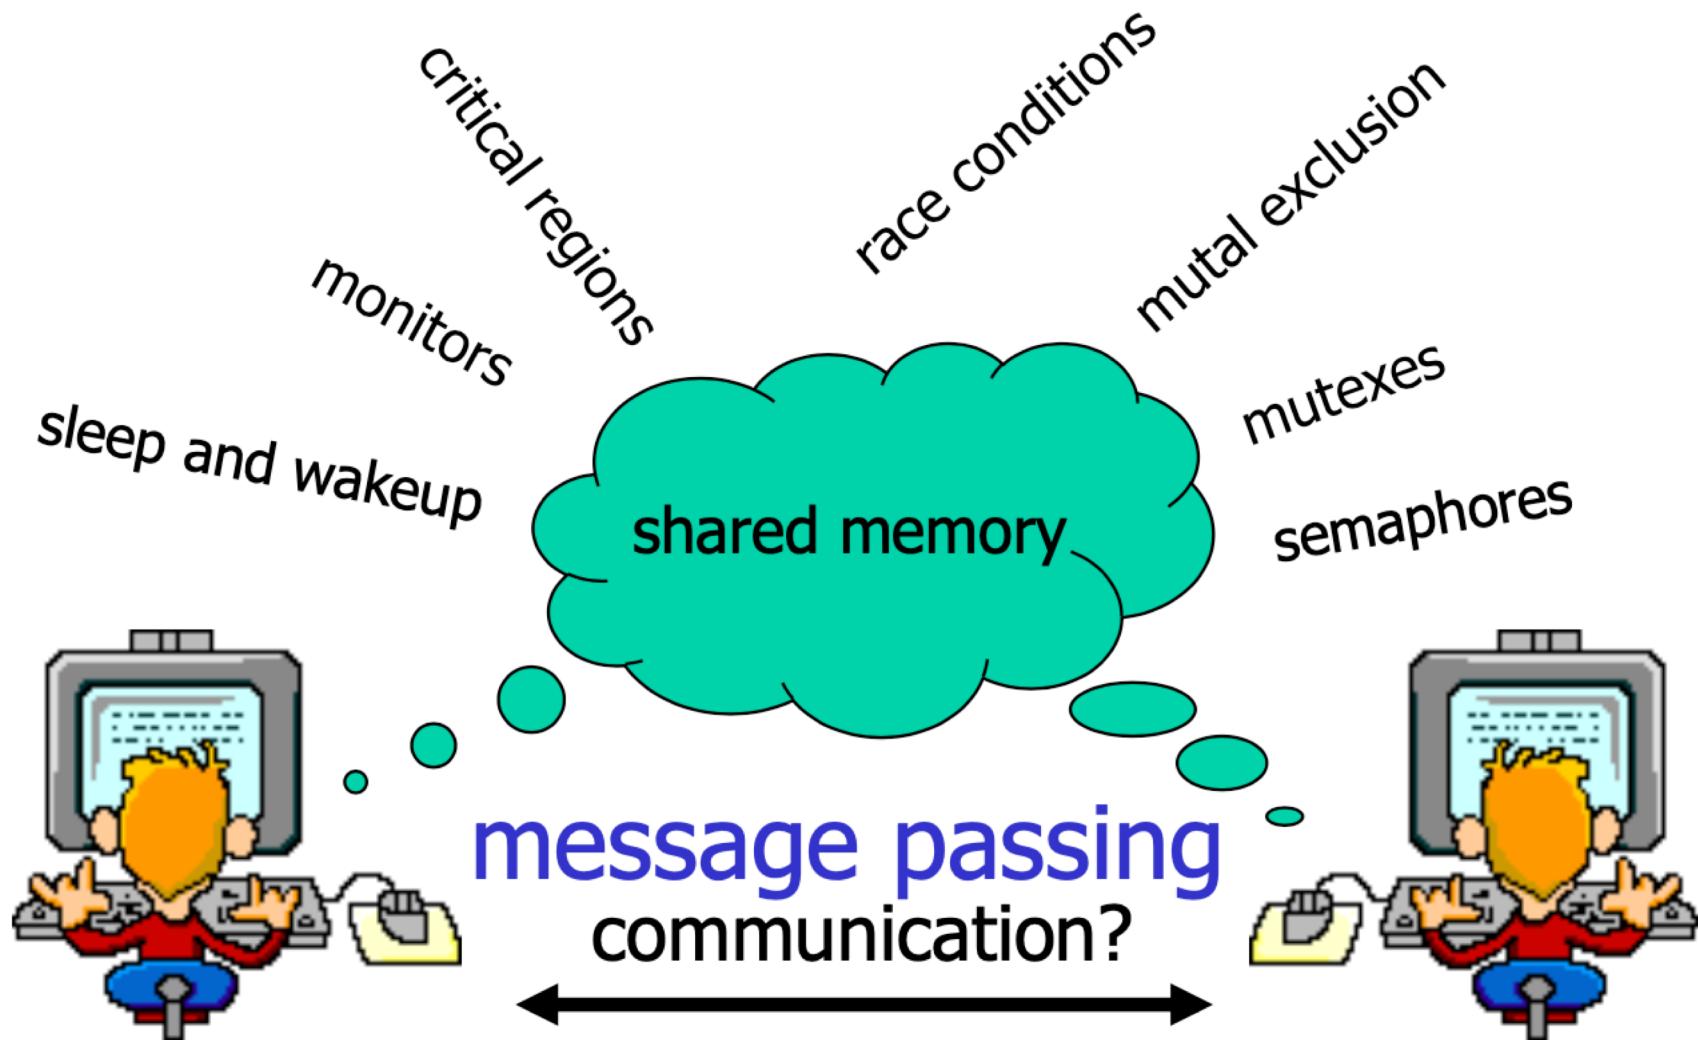
# Concurrent Execution of Processes / Threads

- Concurrency is inherent feature of multiprogramming, multiprocessing, distributed systems, and client-server model of computation
- Implementation of systems using concurrent Threads/Processes require
  - Communication between Processes/Threads
  - Synchronization among processes/Threads
  - Sharing of common resources

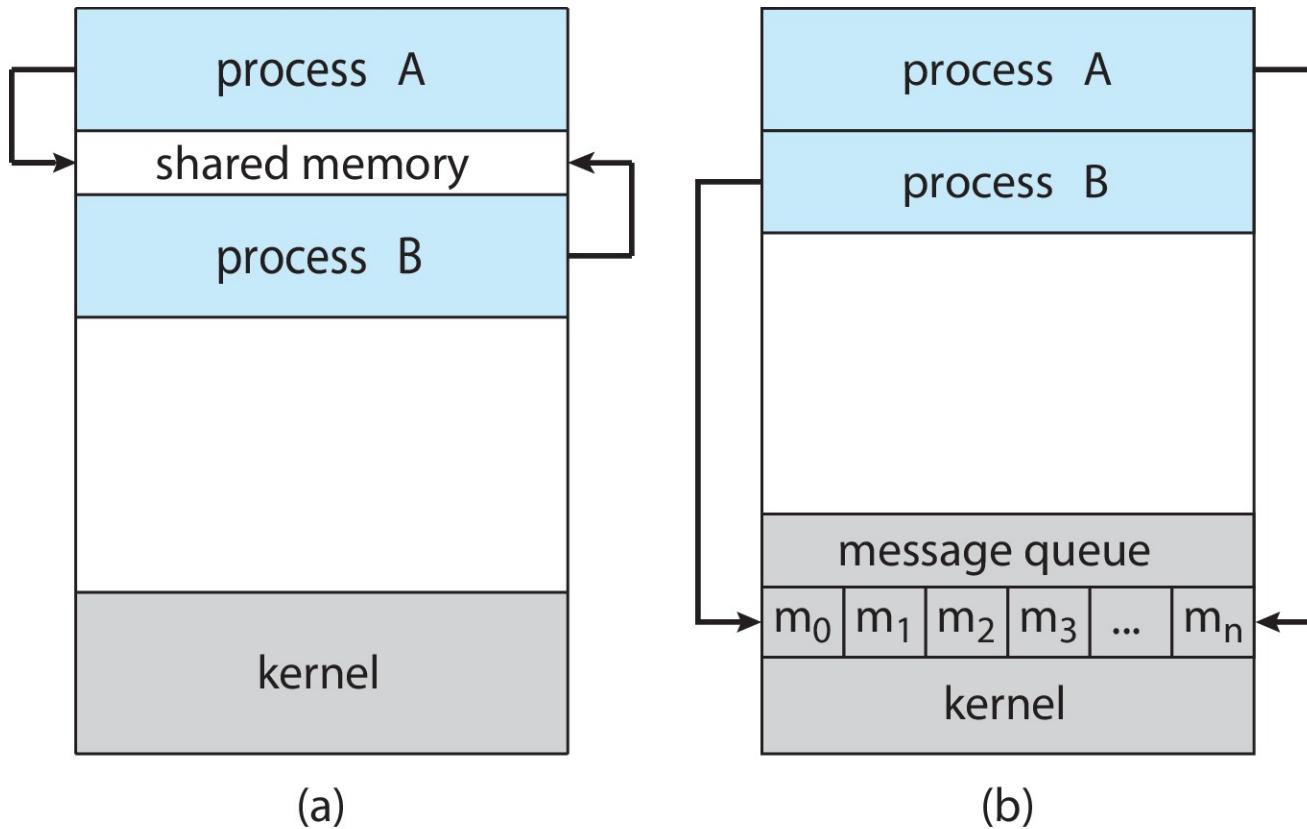
# Interactions among processes

- **Competing processes:** Processes themselves do not share anything. But OS has to share the system resources among these processes “competing” for system resources such as disk, file or printer.
- **Cooperating processes :** Results of one or more processes may be needed for another process. Reasons for cooperation:
  - Information sharing
  - Computation Speedup
  - Modularity
  - Convenience

# IPC – Big Picture



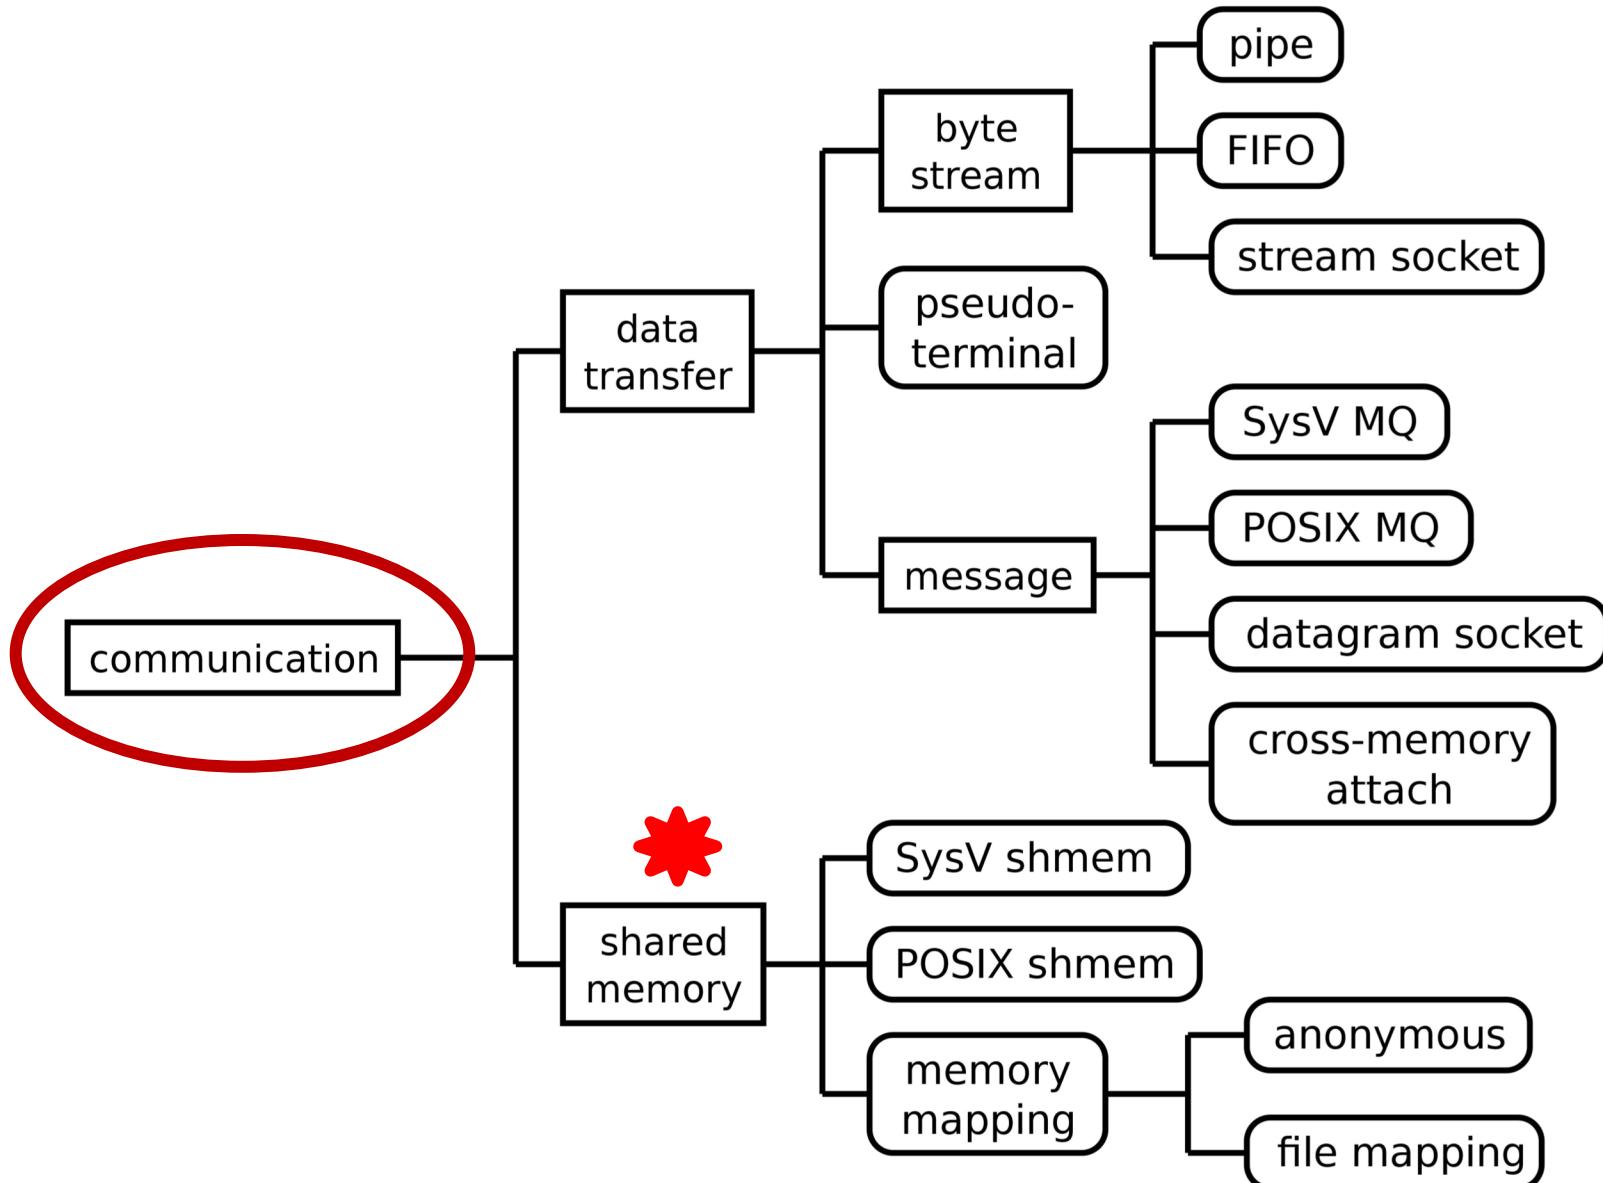
# Models of IPC: Shared Memory, Message Passing



**Shared memory**

**Message passing.**

# Communication



# Models of IPC: Shared Memory, Message Passing

Both models are commonly used in operating systems

- **Shared Memory**

- Faster than message passing as message passing is implemented using system calls (generally slow)
- System Calls are needed only to establish shared memory regions
  - Once shared memory is created, access is the same as routine memory access – thus no help from OS is required

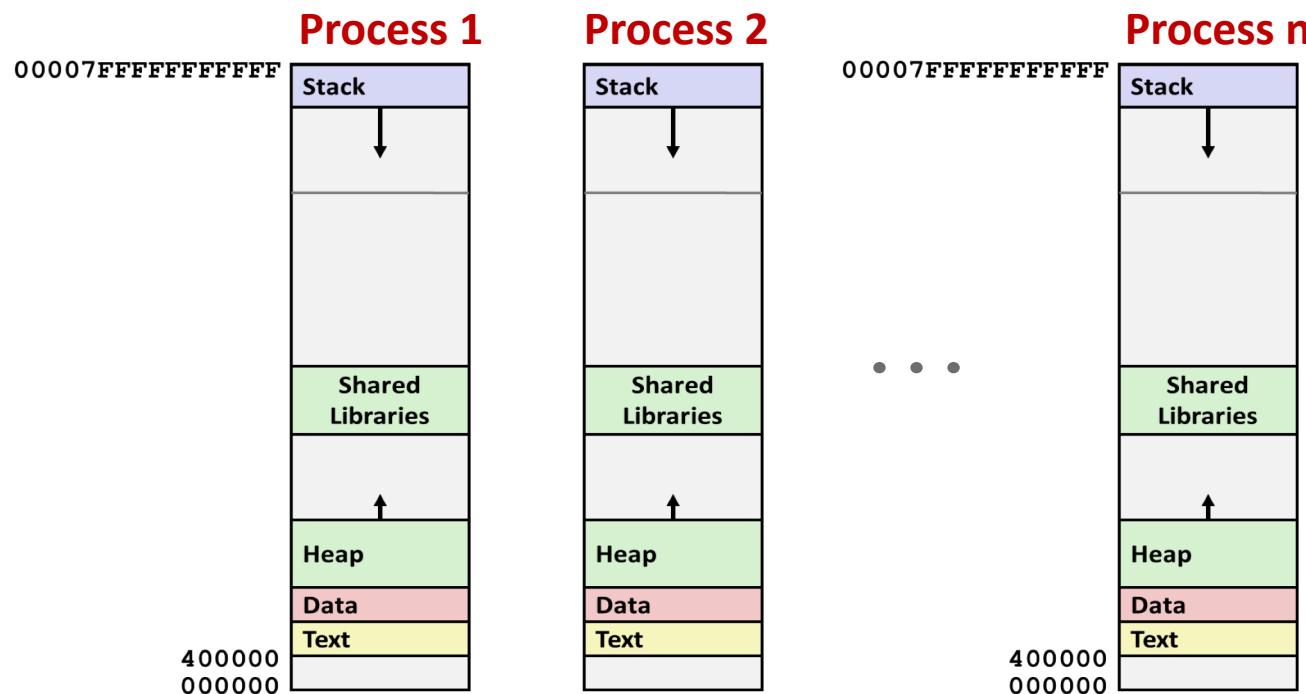
- **Message passing**

- Useful for exchanging small amount of data
- Easier to implement in distributed systems
- Usually implemented using system calls

# **Shared Memory Concept**

# Shared Memory Concept

- Requires communicating processes to establish a region of shared memory
- How to create shared memory region – when each process has its own address space?

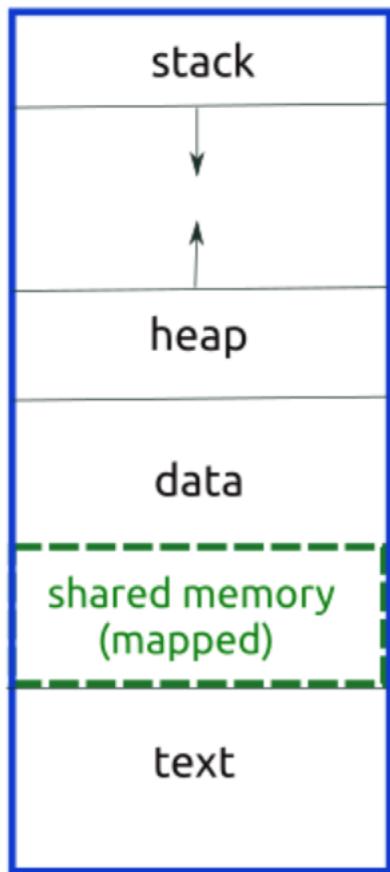


# Shared Memory Concept

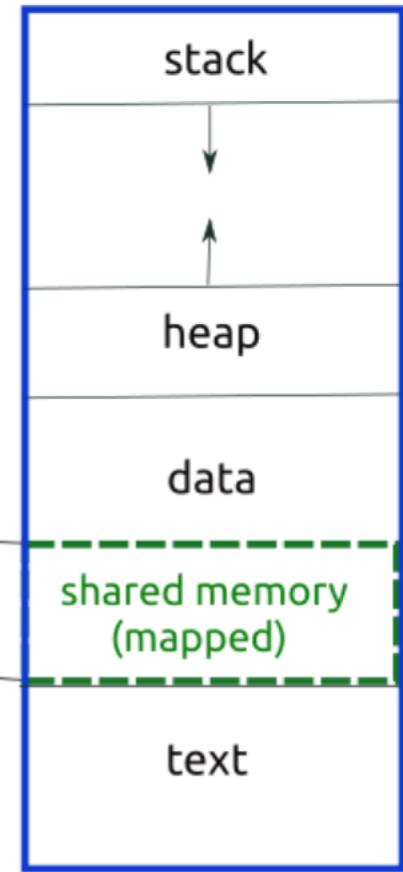
- A shared-memory region resides in the address space of the process creating the shared-memory segment
- Any process that wish to communicate using this shared memory segment must attach it to their address space
- Shared memory requires that two or more processes agree to remove the restriction of non-sharing in their address space
- They can exchange information by reading and writing data into the shared areas
  - Data and location are determined by processes
    - Not under control of OS
- Processes are responsible for ensuring that they are not writing the same location simultaneously
  - Producer Consumer Algorithm

# Shared Memory Model - Steps

- Initiate shared memory mechanism with the help of kernel
- Kernel identifies “safe area” that is attached to both processes
- Attached shared area is used in consistent manner
- When finished, shared data area is detached from all processes which it was attached
- Information regarding the shared memory is deleted from the kernel
- Note that
  - shared memory segment is created by the kernel and mapped to the data segment of the address space of a requesting processes
  - A process can use the shared memory just like any other global variable in its address space



## Shared Memory



# Shared Memory – Fast access

- Processes can read and write shared memory region (segment) directly as ordinary memory access (like they are accessing memory variables directly without kernel help)
  - During this time, kernel is not involved.
  - It is fast – there is no Kernel access
- Race condition can occur

# Procedure for using Shared Memory

- A system wide key is used for identifying shared memory segments
  - Find this key (how?)
  - Use ftok() system call to generate key using a file path and bits of project id
    - This info is also known to other process so they can also generate the key
- Use shmget() to allocate a shared memory
  - Get shared memory id (shmid) as return value
- Use shmat() to attach a shared memory to an address space
  - Use shmid as one of the parameter
  - Any process that has shmid can now attach
- Once communication is over
- Use shmdt() to detach a shared memory from an address space
- Use shmctl() to deallocate a shared memory

# Shared Memory: key generation (1/3)

- To use shared memory, include the following:
  - #include <sys/types.h>
  - #include <sys/ipc.h>
  - #include <sys/shm.h>
- A key is a value of type key\_t.
- There are three ways to generate a key:
  - Do it yourself
  - Use function ftok()
  - Ask the system to provide a private key

# Shared Memory: key generation (2/3)

- Do it yourself: use

```
key_t SomeKey;  
SomeKey = 1234;
```

- Use ftok() to generate one for you:
  - `key_t = ftok(char *path, int ID);`
    - path is a path name (e.g., “.”)
    - ID is an integer (e.g., ‘a’), or Project Id bits
  - Function ftok() returns a key of type `key_t`:
  - `SomeKey = ftok("./", 'x');`
- Keys are global entities. If other processes know your key, they can access your shared memory.
- Ask the system to provide a private key using `IPC_PRIVATE`

# Shared Memory: key generation (3/3)

- First step is to setup shared memory setup in kernel
- Key is generated using ftok() function

```
#include <sys/types.h>
#include <sys/ipc.h>

key_t ftok (const char *pathname, int proj_id);
```

- The *pathname* is an existing file in the filesystem
- The last eight bits of *proj\_id* are used; these must not be zero
- A System V IPC key is returned

# Asking for shared memory (1/4)

- Include the following:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

- Use `shmget()` to request a shared memory:

- `shm_id = shmget(`  
 `key_t key, /* identity key */`  
 `int size, /* memory size */`  
 `int flag); /* creation or use */`

- `shmget()` returns a shared memory ID.
- The flag, for our purpose, is either `0666 (rw)` or  
`IPC_CREAT | 0666`. Yes, `IPC_CREAT`

# Asking for shared memory (2/4)

- The following creates a shared memory of size

```
struct Data with a private key  
IPC_PRIVATE. This is a creation  
(IPC_CREAT) and permits read and write  
(0666).
```

```
struct Data { int a; double b; char x; };  
int ShmID;  
  
ShmID = shmget(  
    IPC_PRIVATE, /* private key */  
    sizeof(struct Data), /* size */  
    IPC_CREAT | 0666); /* cr & rw */
```

# Asking for shared memory (3/4)

- The following creates a shared memory with a key based on the current directory

```
struct Data { int a; double b; char x;};

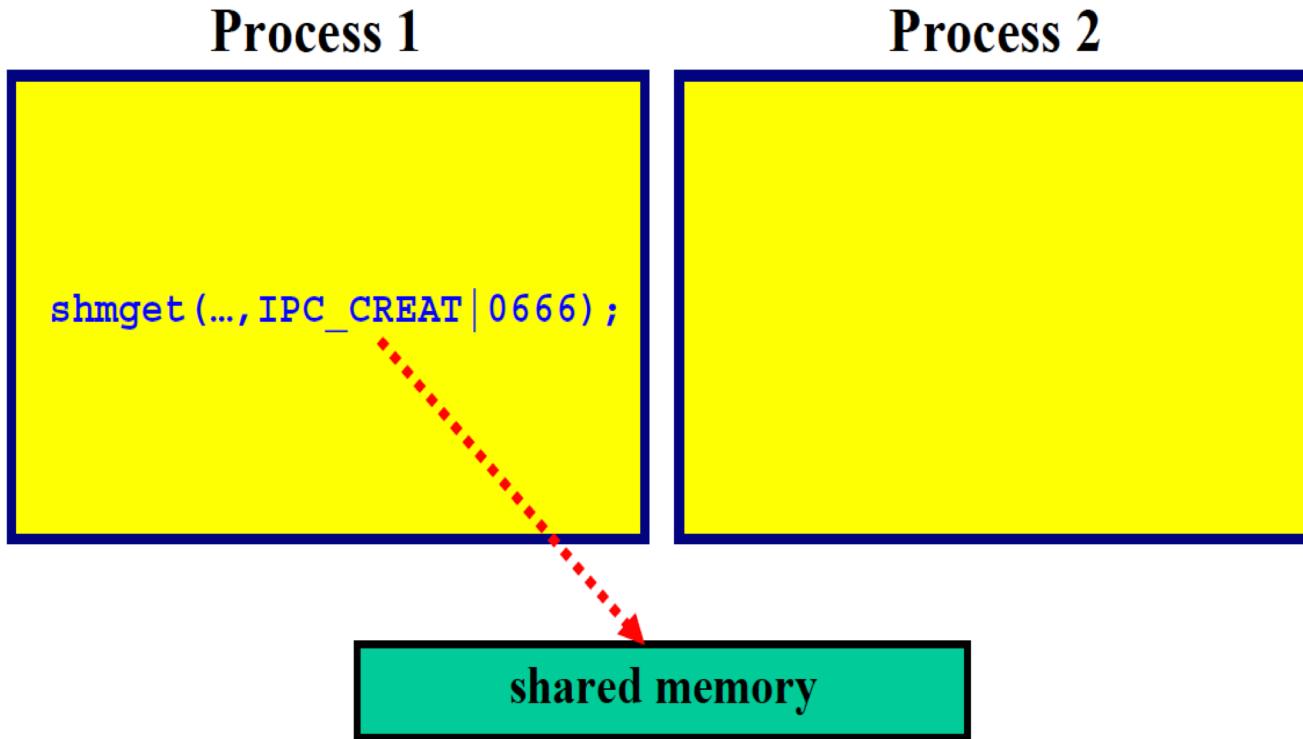
int     ShmID;
key_t   Key;

Key = ftok("./", 'h');
ShmID = shmget(
            Key,          /* a key */
            sizeof(struct Data),
            IPC_CREAT | 0666);
{
```

# Asking for shared memory (4/4)

- The process that creates shared memory
  - It uses flag IPC\_CREAT | 0666
- Other processes that access a created one shared segment using 0666
- If the return value is negative (Unix convention), the request was unsuccessful, and no shared memory is allocated

## After the Execution of `shmget()`



*Shared memory is allocated; but, is not part of the address space*

# Attaching Shared Memory (1/2)

- Use shmat() to attach an existing shared memory to an address space:

```
shm_ptr = shmat(  
    int shm_id,           /* ID from shmget() */  
    char *ptr,            /* use NULL here */  
    int flag);            /* use 0 here */
```

- `shm_id` is the shared memory ID returned by `shmget()`
- Use `NULL` and `0` for the second and third arguments, respectively
- `shmat()` returns a void pointer to the memory.
- If unsuccessful, it returns a negative integer

## Process 1

```
Shmget(..., IPC_CREAT | 0666);  
ptr = shmat(.....);
```

ptr

## Process 2

```
ptr = shmat(.....);
```

ptr

shared memory

*Now processes can access the shared memory*

# Detaching or removing the shared memory

- To detach a shared memory, use

```
shmdt(shm_ptr);
```

`shm_ptr` is the pointer returned by `shmat()`.

- After a shared memory is detached, it is still there. You can re-attach and use it again
- To remove a shared memory, use

```
shmctl(shm_ID, IPC_RMID, NULL);
```

where `shm_ID` is the shared memory ID returned by `shmget()`.

- After a shared memory is removed, it no longer exist

# **Garbage Collection Concept**

# Where does it happen in OS?

# **Do we need to do it for ourselves ?**

# Lecture Summary

- Shared memory provides a fast IPC mechanism