

CS310 Operating Systems

Lecture 20: Need for Synchronization – 2 Too much Milk! and Lock

Ravi Mittal
IIT Goa

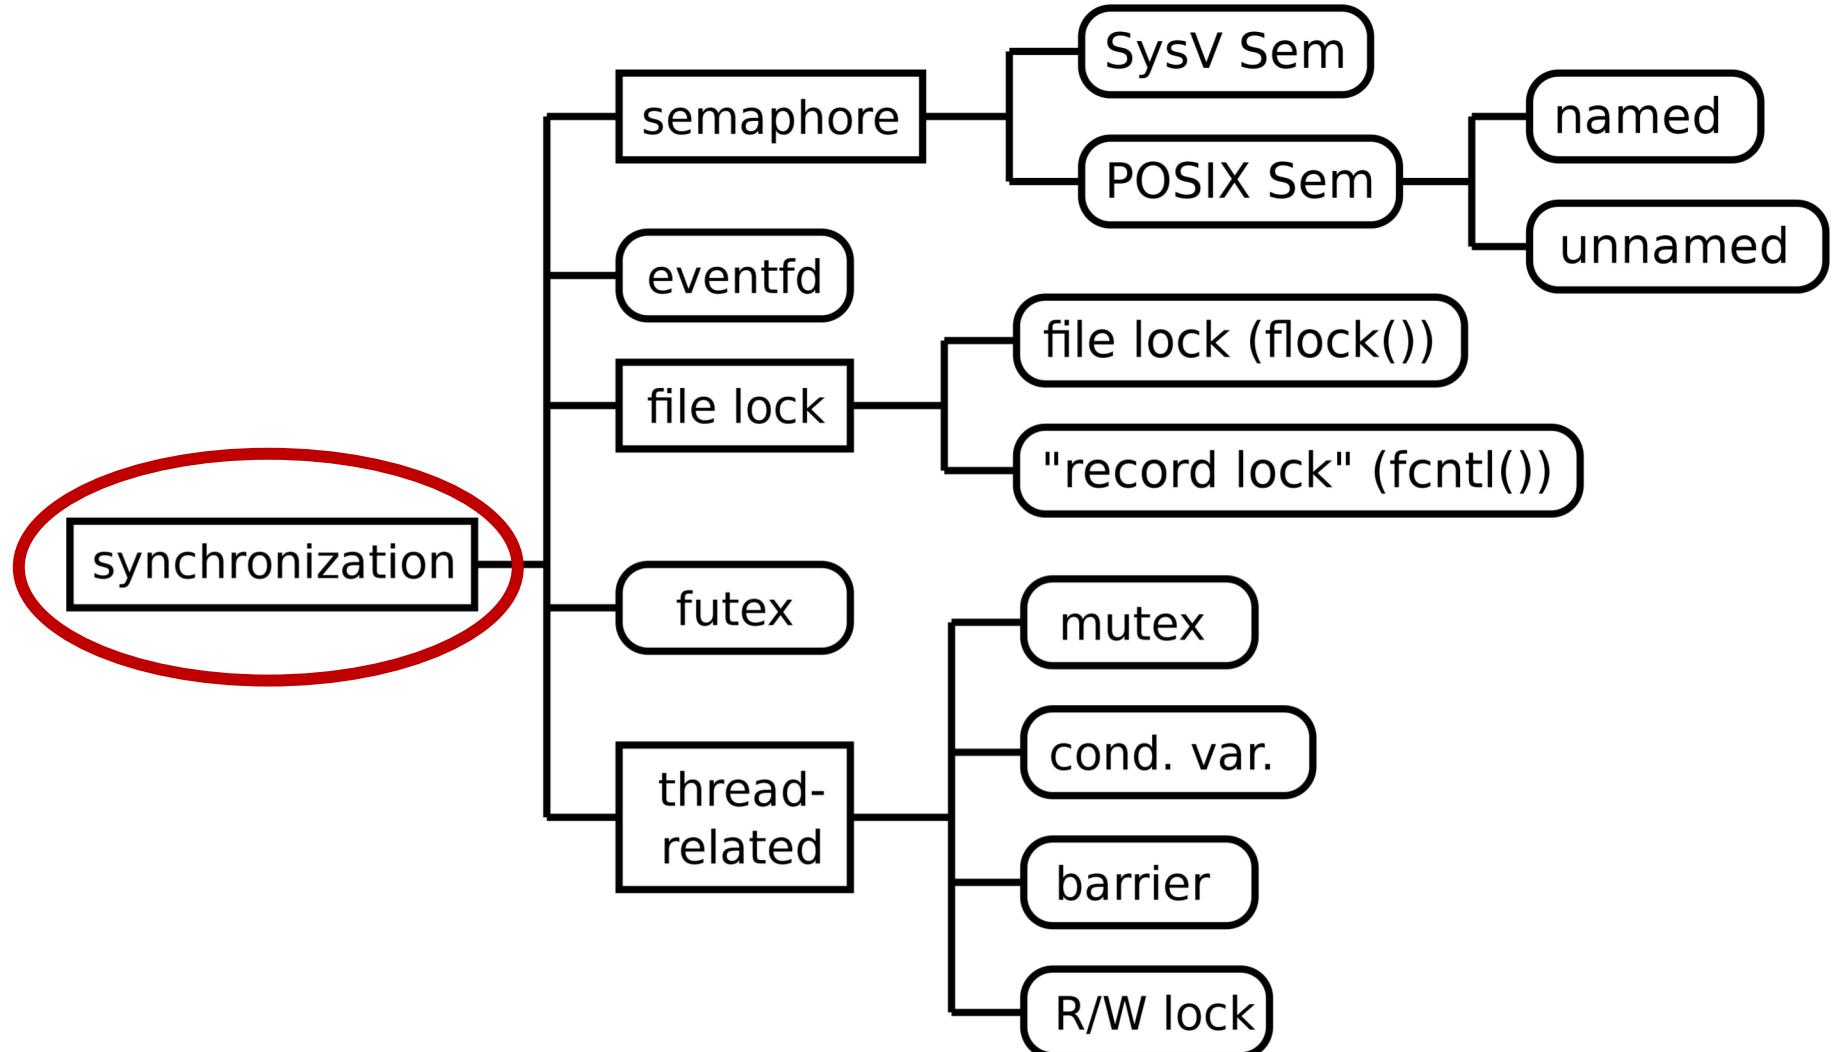
So far we have studied

- Threads
- Processes
- Concurrent execution of Threads and Processes require
 - Communication
 - Synchronization
- In the last class we looked at two cases of race condition with
 - Concurrent Processes
 - Concurrent Threads

Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - CS162, Operating System and Systems Programming, University of California, Berkeley
 - Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Volume II, Chapter 5
 - Section 5.1

Needs for Synchronization



We will start with High level primitives

Programs	Shared Programs			
Higher-level API	Locks Semaphores Monitors Others			
Hardware	Disable Ints Test&Set Compare&Swap, others			

- Our focus will be on concepts

Today we will study ..

- Example: Too much Milk
- Lock definition and properties
- Pthread mutex
- Example: Counting with two threads

Too much Milk !

Motivating Example: “Too Much Milk”

- Great thing about OS’s – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

Use lock to fix Milk problem

- Fix the milk problem by putting a key on the refrigerator
 - Lock it and take key if you are going to go buy milk
 - Fixes too much: roommate angry if only wants Beer



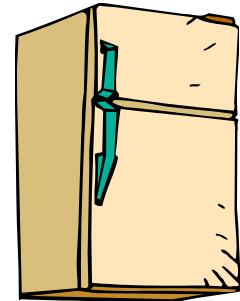
Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
- What are the correctness properties for the “Too much milk” problem???
 - Never more than one person buys
 - Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

Too Much Milk: Solution #1

(1/3)

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this
 - Remember, only memory read/write are atomic



```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```

Does it work properly?

Too Much Milk: Solution #1

(2/3)

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

Thread A

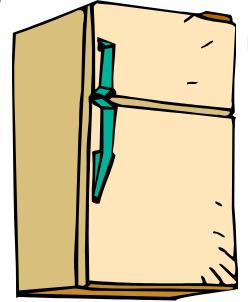
```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy Milk;  
        remove Note;  
    }  
}
```

Thread B

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy Milk;  
        remove Note;  
    }  
}
```

Too Much Milk: Solution #1

(3/3)



- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```

- Result?
 - Still too much milk **but only occasionally!**
 - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...
 - Must work despite what the dispatcher does!

Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (noMilk) {  
    if (noNote) {  
        buy milk;  
    }  
}  
remove Note;
```



- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys milk

Too Much Milk: Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

Thread A

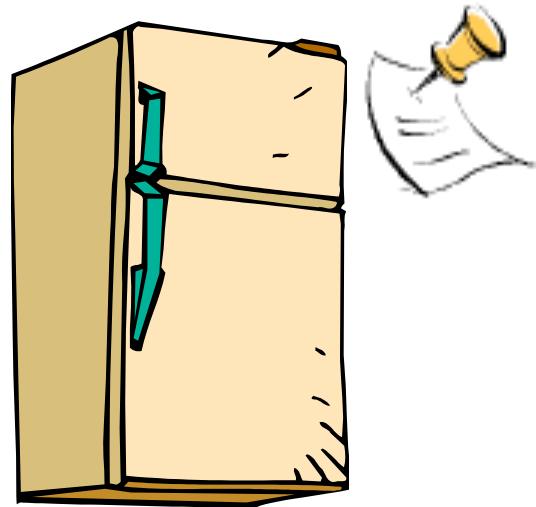
```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNoteA) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note B;
```

- Does this work?
- Possible for neither thread to buy milk
 - Context switches at exactly the wrong times can lead each to think that the other is going to buy

Too Much Milk: Solution #2: problem!



- I'm not getting milk, You're getting milk
- This kind of lockup is called starvation!

Too Much Milk Solution #3 (self reading)

- Here is a possible two-note solution:

Thread A

```
leave note A;  
while (note B) {\\"X  
    do nothing;  
}  
if (noMilk) {  
    buy milk;  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

- Does this work? Yes. Both can guarantee that:

- It is safe to buy, or
 - Other will buy, ok to quit

- At X:

- If no note B, safe for A to buy,
 - Otherwise wait to find out what will happen

- At Y:

- If no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

Solution #3 discussion (self reading)

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex
 - Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - This is called “busy-waiting”
- There's a better way
 - Have hardware provide higher-level primitives than atomic load & store
 - Build even higher-level programming abstractions on this hardware support

Too Much Milk: Solution #4

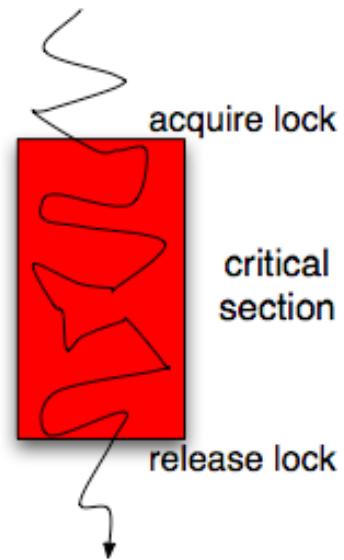
- Suppose we have some sort of implementation of a lock
 - `lock.Acquire()` – wait until lock is free, then grab
 - `lock.Release()` – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
 - Then, our milk problem is easy:

```
milklock.Acquire();  
if (nomilk)    }  
    buy milk;  
milklock.Release();
```

Lock Definition and use

Lock Definition

- Best mainstream solution: **Locks**
 - Implements **mutual exclusion**
 - You can't have it if I have it, I can't have it if you have it
- Two states of a lock: BUSY or FREE
- A lock is initially in the FREE state
- **lock.Acquire()**
 - Waits until the lock is FREE and then atomically makes the lock BUSY
 - Checking the state to see if it is FREE and setting the state to BUSY are together an atomic operation
 - Multiple threads try to acquire the lock, at most one thread will succeed
 - One thread observes that the lock is FREE and sets it to BUSY; the other threads just see that the lock is BUSY and wait



Lock Definition

- `lock.Release()`
 - makes the lock FREE
 - If there are pending acquire operations, this state change causes one of them to proceed
-
- *Only the lock “owner” should release the lock*
- Different OSs have different names for these operations of acquiring and releasing locks
 - `lock()`
 - `Unlock()`

Lock Properties

- Mutual Exclusion (safety property)
 - At most one thread holds the lock
 - Only one thread in the critical section
- Progress (liveness property)
 - If no thread holds the lock and any thread attempts to acquire the lock, then eventually some thread succeeds in acquiring the lock
- Bounded Waiting (liveness property)
 - If thread T attempts to acquire a lock, then there exists a bound on the number of times other threads can successfully acquire the lock before T does

Lock Non-Property: Thread Ordering

- The **bounded waiting property** guarantees that a thread will eventually get a chance to acquire the lock
- However, there is no guarantee that waiting threads acquire the lock in FIFO order

Lock Use

- Consider an update of shared variable
 - $\text{balance} = \text{balance} + 1;$
- This statement forms critical section. We need to protect it with a special lock variable (**mutex** in the example below)

```
1  lock_t mutex; // some globally-allocated lock 'mutex'  
2  ...  
3  lock (&mutex);  
4  balance = balance + 1;  
5  unlock (&mutex);
```

- Note that lock is a variable that holds the state of the lock at any instant in time
- All threads accessing a critical section share a lock
- Only one thread becomes successful in holding the lock – current owner of the lock
 - Thus the thread is in critical section

Lock Use

- `lock()` and `unlock()` are routines
- The thread that acquires the lock enters the critical section
- Then, if another thread calls the lock on the **same lock variable**, it will **not return**
 - So the thread is prevented from entering critical section
 - Function `lock()` doesn't return means it doesn't come out of `lock()` function
- Note that the owner of lock calls unlock
- Thread entities are created by the programmer but scheduled by the OS
 - Locks give some control to the programmer
- `pthreads` library in Linux provides such locks

Pthread lock calls

Pthread lock calls

- **pthread_mutex_init**
 - Creates a new lock in the **unlocked** state
- **pthread_mutex_lock**
 - When the lock is **unlocked**, change the lock to the **locked** state and advance to the next line of code.
 - When the lock is **locked**, this function **blocks** execution of other threads until the lock can be acquired
- **pthread_mutex_unlock**
 - Moves the lock to the **unlocked** state
- Let's use these APIs in our example.

OS Library Locks: *pthreads*

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
                      const pthread_mutexattr_t *attr)  
  
int pthread_mutex_lock(pthread_mutex_t *mutex);  
  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

Mutex Variables

- A typical sequence in the use of a mutex
 1. Create and initialize **mutex**
 2. Several threads attempt to lock **mutex**
 3. Only one succeeds and now owns **mutex**
 4. The owner performs some set of actions
 5. The owner unlocks **mutex**
 6. Another thread acquires **mutex** and repeats the process
 7. Finally **mutex** is destroyed

Example

thread3.c

```
5 int ct = 0;
6
7 void *thread_start(void *ptr) {
8     int countTo = *((int *)ptr);
9
10    int i;
11    for (i = 0; i < countTo; i++) {
12        ct = ct + 1;
13    }
14
15    return NULL;
16}
17
18 int main(int argc, char *argv[]) {
19     // Parse Command Line:
20     if (argc != 3) {
21         printf("Usage: %s <countTo> <thread count>\n",
22               argv[0]);
23         return 1;
24     }
25
26     const int countTo = atoi(argv[1]);
27     if (countTo == 0) { printf("Valid `countTo` is
28 required.\n"); return 1; }
29
30     const int thread_ct = atoi(argv[2]);
31     if (thread_ct == 0) { printf("Valid thread count is
32 required.\n"); return 1; }
33
34     // Create threads:
35     int i;
36     pthread_t tid[thread_ct];
37     for (i = 0; i < thread_ct; i++) {
38         pthread_create(&tid[i], NULL,
39                         thread_start, (void *)&countTo);
40     }
41
42     // Join threads:
43     for (i = 0; i < thread_ct; i++) {
44         pthread_join(tid[i], NULL);
45     }
46
47     // Display result:
48     printf("Final Result: %d\n", ct);
49     return 0;
50 }
```

thread3.c - Observations

- Command line inputs
 - countTo : Each thread counts upto CountTo
 - thread_count: number of threads
- For example:
 - countTo = 1000
 - thread_count = 2
 - → final result = 2000
- However, we see that we don't get exact result for all values of countTo and thread_count
 - At architecture level `ct = ct + 1` is not atomic
 - For example for the first thread – after a few counts – it is context switched after a few steps:
 - Lw s1, ct (assuming ct is in a memory location)
 - Addi s1, 1
 - Context switched.... Other threads now change ct variable
 - After some time the first thread completes operation
 - Ldw s1, ct

Synchronization using mutex

- Mutex : mutual exclusive (to each thread)
- The simplest way to protect a region of code from being accessed is through the use of a **mutex** lock
 - **mutex** stands for *mutual exclusive*

thread3.c

```
5 int ct = 0;
6
7 void *thread_start(void *ptr) {
8     int countTo = *((int *)ptr);
9
10    int i;
11    for (i = 0; i < countTo; i++) {
12        ct = ct + 1;
13    }
14
15    return NULL;
16 }
17
18 int main(int argc, char *argv[]) {
19     // Parse Command Line:
20     if (argc != 3) {
21         printf("Usage: %s <countTo> <thread count>\n",
22               argv[0]);
23         return 1;
24     }
25
26     const int countTo = atoi(argv[1]);
27     if (countTo == 0) { printf("Valid `countTo` is
28 required.\n"); return 1; }
29
30     const int thread_ct = atoi(argv[2]);
31     if (thread_ct == 0) { printf("Valid thread count is
32 required.\n"); return 1; }
33
34     // Create threads:
35     int i;
36     pthread_t tid[thread_ct];
37     for (i = 0; i < thread_ct; i++) {
38         pthread_create(&tid[i], NULL,
39                         thread_start, (void *)&countTo);
40     }
41
42     // Join threads:
43     for (i = 0; i < thread_ct; i++) {
44         pthread_join(tid[i], NULL);
45     }
46
47     // Display result:
48     printf("Final Result: %d\n", ct);
49     return 0;
50 }
```

Non-atomic operation

Critical section – that should be accessed by only one thread at a time

thread3d.c

```
5  pthread_mutex_t lock;
6  int ct = 0;
7
8  void *thread_start(void *ptr) {
9      int countTo = *((int *)ptr);
10
11     int i;
12     for (i = 0; i < countTo; i++) {
13         pthread_mutex_lock(&lock);
14         ct = ct + 1;
15         pthread_mutex_unlock(&lock);
16     }
17
18     return NULL;
19 }
20
21 int main(int argc, char *argv[]) {
22     // Parse Command Line:
23     if (argc != 3) {
24         printf("Usage: %s <countTo> <thread count>\n",
25               argv[0]);
26         return 1;
27     }
28
29     const int countTo = atoi(argv[1]);
30     if (countTo == 0) { printf("Valid `countTo` is
31 required.\n"); return 1; }
32
33     const int thread_ct = atoi(argv[2]);
34     if (thread_ct == 0) { printf("Valid thread count is
35 required.\n"); return 1; }
36
37     // Create Lock:
38     pthread_mutex_init(&lock, NULL);
```

1 Declare lock variable

3 Locking and unlocking
CriticalSection

2 initialize lock to null (open)

Summary

- Concurrent threads are a very useful abstraction
- Concurrent threads introduce problems when accessing shared data
 - Programs must be insensitive to arbitrary interleavings
 - Without careful design, shared variables can become completely inconsistent
- Important concept: Atomic Operations
 - An operation that runs to completion or not at all
 - These are the primitives on which to construct various synchronization primitives
- We will study several mechanisms of implementing locks

Backup Slides

pthread lock calls

- The name that the POSIX library uses for a lock is a **mutex**
 - [`pthread_mutex_init\(\)`](#) - get a mutex
 - [`pthread_mutex_lock\(\)`](#) - lock a mutex (acquire it), block until available
 - [`pthread_mutex_trylock\(\)`](#) - try to lock a mutex (acquire it), do not block if unavailable
 - [`pthread_mutex_unlock\(\)`](#) - unlock a mutex (release it)
 - [`pthread_mutex_destroy\(\)`](#) - destroy a mutex (remove it)

The name that the POSIX library uses for a lock is a mutex,

Pthread Synchronization

- Two primitives
 - Mutex
 - Semaphore with maximum value 1
 - Condition variable
 - Provides a shared signal
 - Combined with a mutex for synchronization

Pthread Mutex

- States
 - Locked
 - Some thread holds the mutex
 - Unlocked
 - No thread holds the mutex
- When several threads compete
 - One wins
 - The rest block
 - Queue of blocked threads

Mutex Variables

- A typical sequence in the use of a mutex
 1. Create and initialize **mutex**
 2. Several threads attempt to lock **mutex**
 3. Only one succeeds and now owns **mutex**
 4. The owner performs some set of actions
 5. The owner unlocks **mutex**
 6. Another thread acquires **mutex** and repeats the process
 7. Finally **mutex** is destroyed

Creating a mutex

```
#include <pthread.h>
int pthread_mutex_init(pthread_mutex_t *mutex,
                      const pthread_mutexattr_t *attr);
```

- Initialize a pthread mutex: the mutex is initially unlocked
- Returns
 - 0 on success
 - Error number on failure
 - **EAGAIN**: The system lacked the necessary resources; **ENOMEM**: Insufficient memory ; **EPERM**: Caller does not have privileges; **EBUSY**: An attempt to re-initialise a mutex; **EINVAL**: The value specified by attr is invalid
- Parameters
 - **mutex**: Target mutex
 - **attr**:
 - NULL: the default mutex attributes are used
 - Non-NULL: initializes with specified attributes

Creating a mutex

- Default attributes
 - Use `PTHREAD_MUTEX_INITIALIZER`
 - Statically allocated
 - Equivalent to dynamic initialization by a call to `pthread_mutex_init()` with parameter `attr` specified as `NULL`
 - No error checks are performed

Destroying a mutex

```
#include <pthread.h>
int pthread_mutex_destroy(pthread_mutex_t *mutex) ;
```

- Destroy a pthread mutex
- Returns
 - 0 on success
 - Error number on failure
 - **EBUSY**: An attempt to re-initialise a mutex; **EINVAL**: The value specified by attr is invalid
- Parameters
 - **mutex**: Target mutex

Locking/unlocking a mutex

```
#include <pthread.h>
int pthread_mutex_lock(pthread_mutex_t *mutex) ;
int pthread_mutex_trylock(pthread_mutex_t *mutex) ;
int pthread_mutex_unlock(pthread_mutex_t *mutex) ;
```

- Returns
 - 0 on success
 - Error number on failure
 - **EBUSY**: already locked; **EINVAL**: Not an initialised mutex; **EDEADLK**: The current thread already owns the mutex; **EPERM**: The current thread does not own the mutex

How to implement Locks?

- How can we build multi-instruction atomic operations?
 - Recall: dispatcher gets control in two ways
 - Internal: Thread does something to relinquish the CPU
 - External: Interrupts cause dispatcher to take CPU
 - On a uniprocessor, can avoid context-switching by:
 - Avoiding internal events
 - Preventing external events by disabling interrupts

Naïve use of Interrupt Enable/Disable

- Naïve Implementation of locks:

```
LockAcquire { disable Ints; }  
LockRelease { enable Ints; }
```

- Problems with this approach:

- Can't let user do this!
 - Consider following:

```
LockAcquire();  
While(TRUE) {} // infinite loop
```

- Critical Sections might be arbitrarily long
- What happens with I/O or other important events?
 - Flight control

backup

Too Much Milk Solution #3 (self reading)

- Here is a possible two-note solution:

Thread A

```
leave note A;  
while (note B) {\\"X  
    do nothing;  
}  
if (noMilk) {  
    buy milk;  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - If no note B, safe for A to buy,
 - Otherwise wait to find out what will happen
- At Y:
 - If no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

Solution #3 discussion (self reading)

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex
 - Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - This is called “busy-waiting”
 - There's a better way
 - Have hardware provide higher-level primitives than atomic load & store
 - Build even higher-level programming abstractions on this hardware support

Case 1

- “leave note A” happens before “if (noNote A)”

The diagram illustrates a dependency between two code snippets. A blue arrow labeled "happened before" points from the first snippet (X) to the second snippet (Y).

```
leave note A;  
while (note B) {\\"X  
    do nothing;  
};  
  
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;  
  
if (noMilk) {  
    buy milk; }  
}  
remove note A;
```

Case 1

- “leave note A” happens before “if (noNote A)”

```
leave note A;  
while (note B) {\\"X  
    do nothing;  
};
```

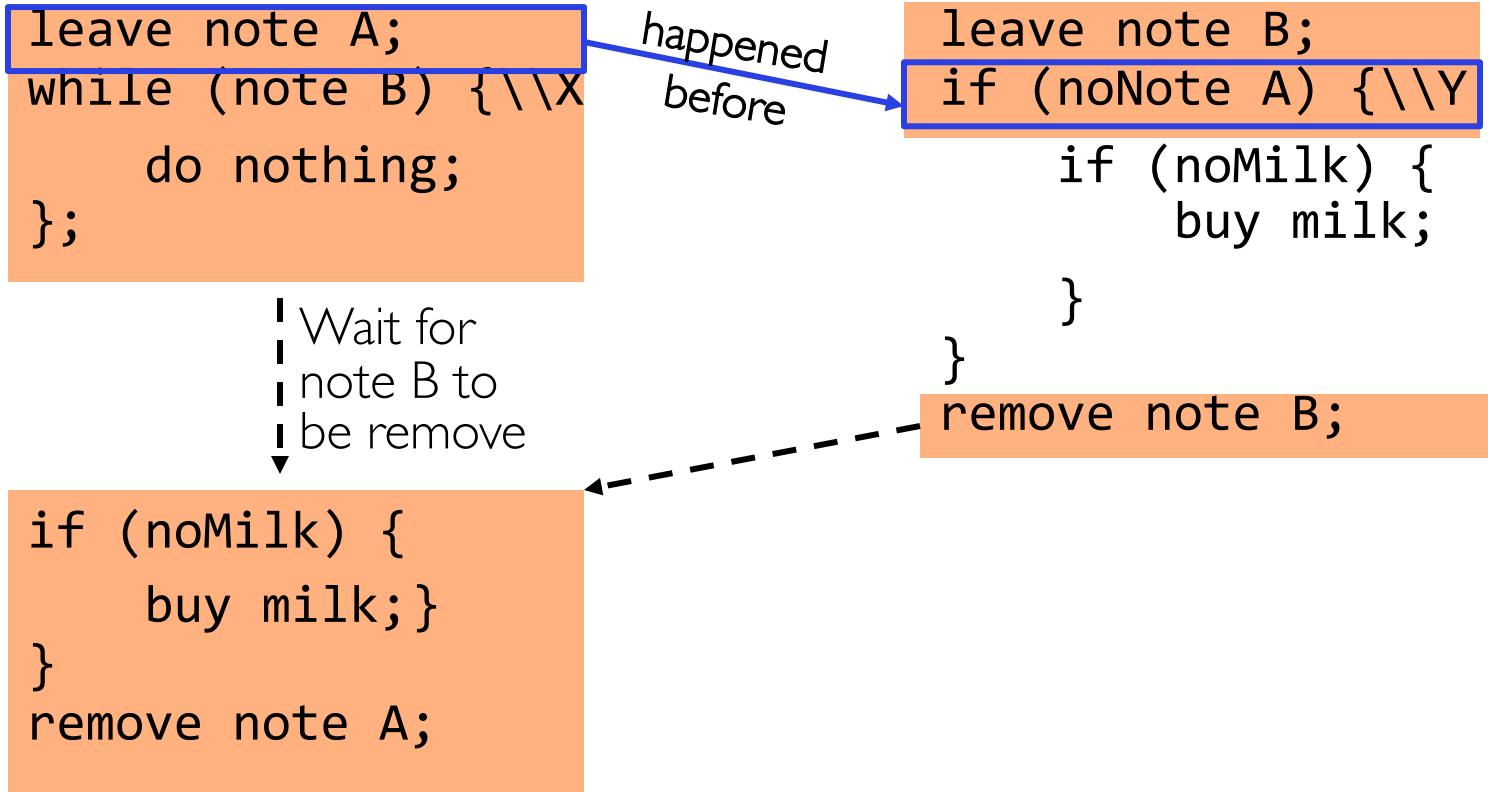
```
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

```
if (noMilk) {  
    buy milk; }  
}  
remove note A;
```

happened
before

Case 1

- “leave note A” happens before “if (noNote A)”



Case 2

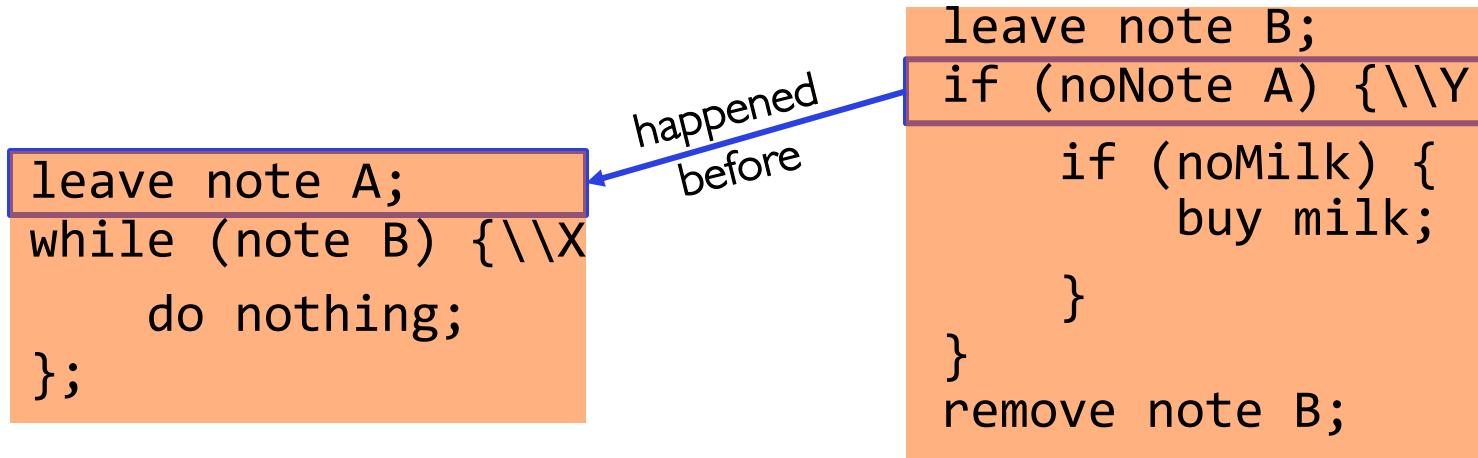
- “if (noNote A)” happens before “leave note A”

```
leave note A;  
while (note B) {\\"X  
    do nothing;  
};  
  
leave note B;  
if (noNote A) {\\"Y  
    if (noMilk) {  
        buy milk;  
    }  
    remove note B;
```

```
if (noMilk) {  
    buy milk; }  
}  
remove note A;
```

Case 2

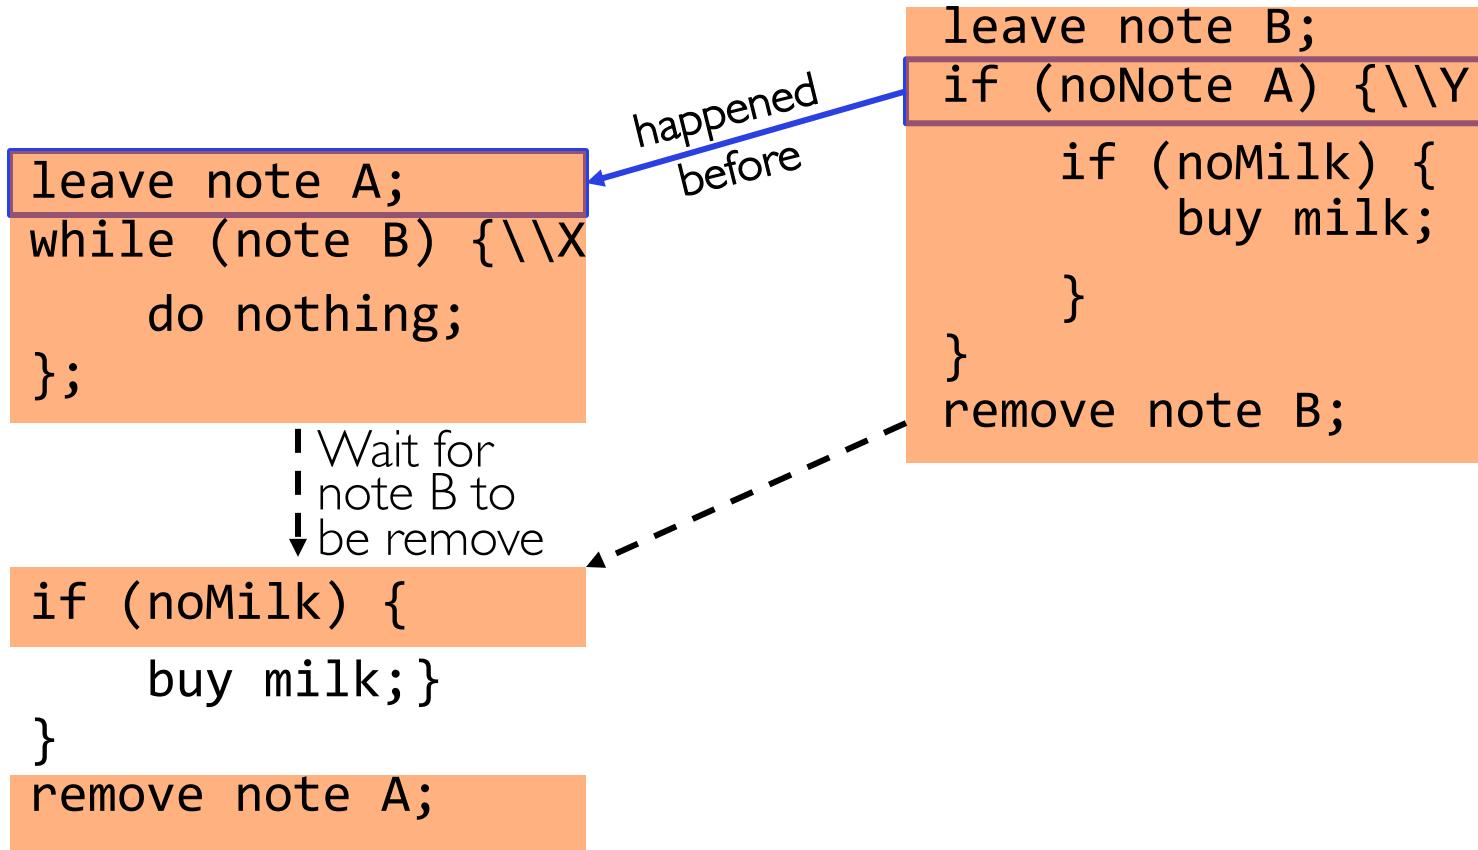
- “if (noNote A)” happens before “leave note A”



```
if (noMilk) {  
    buy milk; }  
}  
remove note A;
```

Case 2

- “if (noNote A)” happens before “leave note A”



Atomic Operations

- Indivisible operations that cannot be interleaved with or split by other operations
- What is the need for atomic operations?