

CS310 Operating Systems

Lecture 21 : Locks using Hardware atomic Instructions: Test-and-Set, Compare-and-swap

Ravi Mittal
IIT Goa

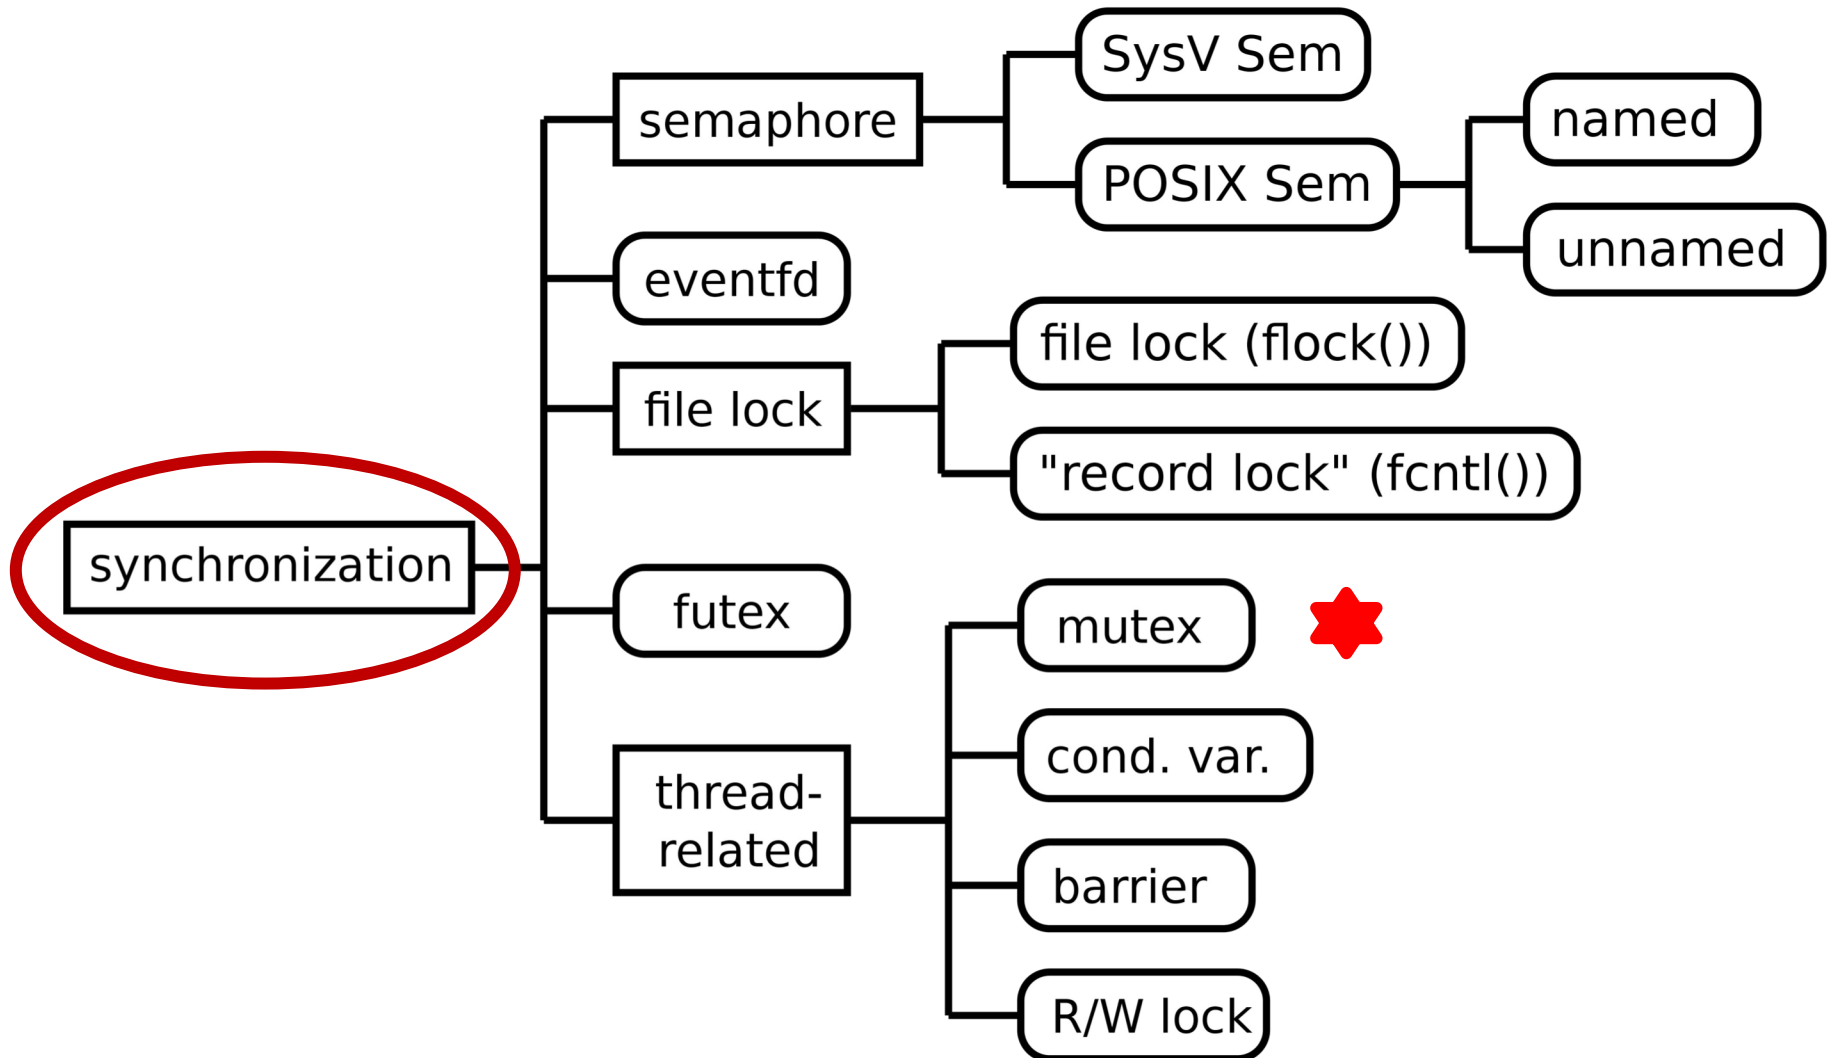
Acknowledgements !

- Contents of this class presentation has been taken from various sources. Thanks are due to the original content creators:
 - Book: Operating System: Three Easy Pieces, by Remzi H Arpaci-Dusseau, Andrea C Arpaci-Dusseau, Chapter 28 Locks
 - <https://pages.cs.wisc.edu/~remzi/OSTEP/threads-locks.pdf>

Reading

- Book: Operating System: Three Easy Pieces, by Remzi H Arpaci-Dusseau, Andrea C Arpaci-Dusseau
 - Chapter 31,
<https://pages.cs.wisc.edu/~remzi/OSTEP/threads-sema.pdf>
- Book: Operating Systems: Principles and Practice: Thomas Anderson and Michael Dahlin, Volume II, Chapter 5.8

Needs for Synchronization



We will start with High level primitives

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors Others
Hardware	Disable Ints Test&Set Compare&Swap, others


In the class we will study

- Implementations of locks
- Lock implementation by disabling interrupts
- Test-and-Set Atomic Instruction
- Compare-and-swap
- Sleeping Mutex
- Using Queues to improve on spinning

Last Class: Lock Introduction

What if we want to generalize beyond increments?

- Best mainstream solution: **Locks**
 - Implements **mutual exclusion**
 - You can't have it if I have it, I can't have it if you have it



when lock = 0, set lock = 1, continue

lock = 0

Recall: Lock

- Consider an update of shared variable
 - `balance = balance + 1;`
- This statement forms critical section. We need to protect it with a special lock variable (mutex in the example below)

```
1  lock_t mutex; // some globally-allocated lock 'mutex'
2  ...
3  lock(&mutex);
4  balance = balance + 1;
5  unlock(&mutex);
```

- Note that lock is a variable that holds the state of the lock at any instant in time
- All threads accessing a critical section share a lock
- Only one thread becomes successful in holding the lock – current owner of the lock
 - Thus the thread is in critical section

Recall: Lock

- `Lock()` and `unlock()` are routines
- The thread that acquires the lock enters the critical section
- Then, if another thread calls the lock on the same lock variable, it will not return
 - So the thread is prevented from entering critical section
 - Function `lock()` doesn't return means it doesn't come out of `lock()` function
- Note that the owner of lock calls `unlock`
- Thread entities are created by the programmer but scheduled by the OS
 - Locks give some control to the programmer
- Pthreads library in Linux provides such locks

pthread lock calls

- The name that the POSIX library uses for a lock is a **mutex**
- [pthread_mutex_init\(\)](#) - get a mutex
- [pthread_mutex_lock\(\)](#) - lock a mutex (acquire it), block until available
- [pthread_mutex_trylock\(\)](#) - try to lock a mutex (acquire it), do not block if unavailable
- [pthread_mutex_unlock\(\)](#) - unlock a mutex (release it)
- [pthread_mutex_destroy\(\)](#) - destroy a mutex (remove it)

Lock Implementation questions

- How can we build an efficient lock?
- What hardware support is required?
- What OS support is required?

Building a lock

- Goals of lock implementation
 - Mutual Exclusion
 - Fairness
 - Performance
 - Low time overhead
 - Acquiring, Releasing and waiting for a lock should not consume too many resources
- Implementations of locks are needed for both user-space programs and kernel code
- Implementing locks needs support from **hardware** and **OS**

Lock implementation by disabling interrupts

- The earliest solution: mutual exclusion by disabling interrupts
- Solution was invented for a single processor system

```
1 void lock() {  
2     DisableInterrupts();  
3 }  
4 void unlock() {  
5     EnableInterrupts();  
6 }
```

- Simple approach
- Is it good enough? No.
- Disabling interrupts is a privileged instruction (for turning interrupts on and off) and program can misuse it (e.g., run forever)
- This approach will not work on multiprocessor systems, since another thread on another core can enter critical section

Lock implementation by disabling interrupts

- Turning off interrupts for extended period of time can lead to other interrupts becoming lost - this is a serious problem to handle
- It is dangerous to give a user process a power to disable interrupt
 - What happens if a user process disable interrupts and goes into infinite loop?
- This approach is inefficient
- This approach of interrupt masking is used by kernel to guarantee atomicity when accessing its own data structures
 - No trust issue inside the OS

Does this lock implementation work?

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 -> lock is available, 1 -> held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1;          // now SET it!
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

- First thread that enters the critical section calls `lock()`
 - Tests if flag is 1; as the flag is 0, it enters the critical section and setup flag = 1
 - Thread 1 now holds the flag
 - When it finishes it calls `unlock()`
- Another thread calls `lock()` when thread 1 is in critical section
- It **spin-waits** in the while loop until flag becomes 0

Does this lock implementation work? **No**

- The solution is incorrect. Consider the code interleaving

Thread 1	Thread 2
call lock ()	
while (flag == 1)	
interrupt: switch to Thread 2	
	call lock ()
	while (flag == 1)
	flag = 1;
	interrupt: switch to Thread 1
flag = 1; // set flag to 1 (too!)	

- Thread 1 spins and finds out lock flag is 0;
- Thread 1 is interrupted before it sets flag = 1
- Thread 2 also finds flag = 0 and sets flag = 1 and enters critical section
- Thread 2 is interrupted
- Thread 1 now sets flag = 1 and enters critical section
- Also, Spin-waiting wastes processor's time --> expensive

Atomic Read-Modify-Write Instructions

- Hardware instructions that allows us to test and set or compare and swap, operations atomically

- Test_and_Set
 - Compare_and_Swap
 - Load-Linked and Store-Conditional
 - Fetch-And-Add
- 
- Will study

- We can build locks with these instruction

Test-and-Set: Hardware Atomic Instruction

- It's hard to ensure atomicity only in software
- Modern architectures provide hardware atomic instructions
 - Test-and-Set
- It enables to test the old value (which is what is returned) while simultaneously setting the memory location to a new value
 - Function returns the old value

Note that operation (of reading old value (of lock variable) and setting the new value) is atomic

This is used to implement lock

Test-and-Set: Hardware Atomic Instruction

```
1      int TestAndSet(int *old_ptr, int new) {  
2          int old = *old_ptr; // fetch old value at old_ptr  
3          *old_ptr = new;      // store 'new' into old_ptr  
4          return old;          // return the old value  
5      }
```

- It returns the old value pointed to by the `old_ptr`, and simultaneously updates said value to new

Test-and-Set: Hardware Atomic Instruction

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0: lock is available, 1: lock is held
7      lock->flag = 0;
8  }
9
10 void lock(lock_t *lock) {
11     while (TestAndSet(&lock->flag, 1) == 1)
12         ; // spin-wait (do nothing)
13 }
14
15 void unlock(lock_t *lock) {
16     lock->flag = 0;
17 }
```

- Case 1: Flag is 0 initially. Thread 1 calls `lock()`.
 - Calls `TestAndSet(flag, 1)`
 - Routine returns old value of flag i.e. 0 and then atomically sets flag to 1; thread enters critical section
 - When the thread is finished it sets the flag to 0
- Case 2: Thread 1 is holding the lock. The second thread calls `lock()` and then calls `TestAndSet()`. `TestAndSet` returns the old value of flag i.e. 1 and simultaneously setting it to 1
 - Thread 2 spin-locks

Evaluating Spin Locks

- Correctness – Yes
 - It provides mutual exclusion
- Fairness
 - Spin-locks don't provide any fairness guarantees
 - An unlucky thread may spin forever under contention
 - Starvation
- Performance
 - Single CPU case
 - Poor performance. If the thread holding lock is preempted, all other threads that are competing for the lock will spin wait in their time slot

Spin lock with **compare-and-swap**

- Another hardware primitive – some systems provide it
- Idea: To test whether the value at the address specified by `ptr` is equal to `expected`
 - if so, update the memory location pointed to by `ptr` with the new value
 - If not, do nothing
 - In both cases, return the original value
 - Thus the code calling `compare-and-swap` come to know whether it succeeded or not
- Similar to `test-and-set` instruction

Spin lock with **compare-and-swap**

- **compare-and-swap** Instruction

```
1  int CompareAndSwap(int *ptr, int expected, int new) {  
2      int original = *ptr;  
3      if (original == expected)  
4          *ptr = new;  
5      return original;  
6  }
```

- Spinlock using compare-and-swap

```
1  void lock(lock_t *lock) {  
2      while (CompareAndSwap(&lock->flag, 0, 1) == 1)  
3          ; // spin  
4  }
```


Simple Approach: Just Yield (sleeping mutex)

- What to do when a context switch occurs in a critical section?
 - Other threads start to spin endlessly, waiting for the interrupted (lock-holding) thread to be run again
- Instead of spinning for a lock, a contending thread simply gives up the CPU and check back later
 - – `yield()` moves thread from **running** to **ready** state

```
1  void init() {
2      flag = 0;
3  }
4
5  void lock() {
6      while (TestAndSet(&flag, 1) == 1)
7          yield(); // give up the CPU
8  }
9
10 void unlock() {
11     flag = 0;
12 }
```

Spin lock and Sleeping Mutex

- If there are many threads (say 100) they contend for lock repeatedly
- If a thread holding lock is preempted, other 99 threads will each call `lock()` and find it held
 - **Sleeping Mutex**: These threads will **yield** as soon it finds that lock is held
 - **Spin lock**: These threads will waste entire time slot spin locking
- Context switch can be substantial → expensive
- Starvation problem still exist

Using Queues – to improve on spinning

- Too much is left to chance
- if the scheduler makes a bad choice, it picks up a thread
 - that is spin-waiting for the lock (1st approach)
 - that yields the CPU immediately (2nd approach)
- A lot of waste and possible starvation
- Can we have more control over which thread is scheduled next?
- Keep a **queue** to track which threads are waiting to acquire the lock – used in Solaris OS
 - Schedule other thread that doesn't want to use **lock()**
- Many such ideas exist

How should locks be used?

- Thread-safe data structures
 - A lock should be acquired before accessing any variable or data structure that is shared between multiple threads
- Coarse-grained Locking
 - One big lock for all shared data
 - Easy to manage
- Fine-grained locking
 - Separate locks
 - More parallelism
 - Difficult to manage
- Note that OS only provides locks, correct locking discipline is left to users

Lecture Summary



- Spin-lock based approaches
 - Test-and-Set: Atomic -> hence lock is atomic
 - Compare-and-Swap: Atomic -> hence lock is atomic
 - Spin-lock based approaches are inefficient
 - **Priority Inversion**: If busy-waiting thread has higher priority than thread holding lock \Rightarrow no progress!
 - Unbounded waiting
- Locks inside the OS are always spinlocks
- Sleeping Mutex Approach
 - Used in most user space lock implementations
- Priority Inversion problem with original Martian rover